

THE SUPPLY OF SKILL AND ENDOGENOUS TECHNICAL CHANGE: EVIDENCE FROM A COLLEGE EXPANSION REFORM

Pedro Carneiro

University College London, Institute for
Fiscal Studies, and Centre for Microdata
Methods and Practice, UK, and FAIR
NHH, Norway

Kai Liu

Faculty of Economics, University of
Cambridge, UK and Department of
Economics, Norwegian School of
Economics, Norway

Kjell G. Salvanes

Norwegian School of Economics, CELE,
and FAIR NHH, Norway

Abstract

We examine the labor market consequences of an exogenous increase in the supply of skilled labor in several municipalities in Norway, resulting from the construction of new colleges in the 1970s. We find that skilled wages increased as a response, suggesting that along with an increase in the supply there was also an increase in demand for skill. We also show that college openings led to an increase in the productivity of skilled labor and investments in R&D. Our findings are consistent with models of endogenous technical change where an abundance of skilled workers may encourage firms to adopt skill-complementary technologies. (JEL: J23, J24, O33, I24)

1. Introduction

There are strong links between technological progress and labor markets. Skill-biased technical change (SBTC) is likely to lead to an increase in the skill premium (see

The editor in charge of this paper was Giovanni Peri.

Acknowledgments: We thank Daron Acemoglu, Gordon Dahl, Magne Mogstad, Robert Moffit, Jarle Møen, and Jeff Smith, as well as the participants at seminars and conferences, for helpful comments and fruitful discussions. This work was supported by the Research Council of Norway through its Centres of Excellence Scheme, FAIR project No. 262675. Pedro Carneiro gratefully acknowledges financial support from the Economic and Social Research Council through the Centre for Microdata Methods and Practice (ES/P008909/1), and support of the European Research Council through ERC-2015-CoG-682349. Kjell G. Salvanes also gratefully acknowledges the support by the Research Council of Norway Grant 202445.

E-mail: p.carneiro@ucl.ac.uk (Carneiro); kai.liu@econ.cam.ac.uk (Liu);
kjell.salvanes@nhh.no (Salvanes)

Journal of the European Economic Association 2023 21(1):48–92

<https://doi.org/10.1093/jeea/jvac032>

© The Author(s) 2022. Published by Oxford University Press on behalf of European Economic Association. This is an Open Access article distributed under the terms of the Creative Commons Attribution License (<http://creativecommons.org/licenses/by/4.0/>), which permits unrestricted reuse, distribution, and reproduction in any medium, provided the original work is properly cited.

e.g. Katz and Murphy 1992; Autor, Katz, and Krueger 1998; Autor, Katz, and Kearney 2008), which, in turn, becomes an incentive for individuals to acquire more skills. At the same time, changes in the supply of skills affect the returns to using skill-complementary technologies and may induce firms to upgrade their technology. The latter mechanism is emphasized in Acemoglu (1998) and Beaudry and Green (2003).

In these papers, an inflow of skilled workers increases returns to using more skill-complementary technologies. If the inflow of skills becomes sufficiently large, then firms upgrade their technology. Initially, the skill premium decreases as we move along a downward-sloping demand curve. Once the increase in the supply is large enough for firms to invest in a new technology, the demand for skill shifts outward. As a result, the skill premium and the supply of skilled workers may increase simultaneously.

Our paper provides new evidence that an exogenous shock to the supply of skilled labor induces endogenous technical change. We study data from a college expansion reform in Norway, which was rolled out across local labor markets and expanded the supply of college-educated workers, and investigate what happened to wages of skilled workers, the productivity of skilled workers, and R&D investments by firms.

We document three main empirical results. First, following the opening of a college, both the relative supply of skilled workers and their relative earnings increase simultaneously. In the years immediately after the reform, the increase in the relative supply of skills occurs mainly among young workers (due to the inflow of new university graduates), whereas the increase in the relative earnings of skilled workers occurs mainly among older workers. In the longer run, a college opening induces increases in the relative supply and earnings of skilled workers who were both young and old at the time of the reform.

These empirical results are consistent with a model where young and old workers are imperfect substitutes (Card and Lemieux 2001). The earnings of older skilled workers are not very much subject to the downward pressure induced by an increase in the supply of skilled workers, and increase shortly after the opening of a college because of endogenous SBTC. The earnings of skilled young workers are also affected by endogenous SBTC but are more subject to downward pressure from the increase in supply. Moreover, these patterns are much more pronounced following the opening of a STEM college than following the opening of a non-STEM college. Increases in the incentive to invest in new technologies occur mainly in areas where there are increases in the supply of skilled workers in STEM fields.¹

1. The Card and Lemieux (2001) model is stylized: It decomposes relative wage changes into contributions from relative supply changes and relative demand changes, where the relative demand changes are labelled as “technology changes”. It is agnostic with regard to the specific channel driving the technical change. More specifically, the technical change identified from the Card and Lemieux framework cannot reveal whether the technical change is driven by any human capital spillover effects or firms using technology that is more complementary to skilled labor. Our firm-level analysis provides more direct

Second, following the opening of a college, both the supply of skilled workers and their marginal productivity increase simultaneously. The marginal product of skilled and unskilled labor is estimated using plant-level production functions, relying on information on output and input factors (and ignoring any wage data).

Third, following the opening of a college, firms invest more in R&D (both in terms of expenditure and employment). Together, these three findings suggest that firms responded to the opening of a college, and the resulting increase in the availability of skilled labor, by promoting technical change, either through the adoption of skill-augmenting technologies or changes in organizational form (Acemoglu 1998; Beaudry and Green 2003).

We argue throughout the paper that, in the period under study, new colleges were not engaged in R&D or innovation activities themselves (we document that no new patents were registered). They were essentially producing new graduates, so their impact on technical change only occurred indirectly, through the endogenous response of firms to an increase in skill supplies. Moreover, we also show that the construction of new colleges is unlikely to have caused substantial increases in the demand for skilled individuals since the employment in the college sector was tiny relative to the size of the labor markets in the locations where new colleges were established. Finally, we explain that our findings are not affected by migration since migration does not respond to college construction, and we are also able to rule out any trade-based explanations of changes in factor prices.

Our empirical analysis is based on several population-wide and long panel data sets, containing rich firm-level information on inputs and outputs, and individual-level information on demographics and labor market outcomes. Firm-level data are available for the population of plants in the manufacturing sector in Norway, spanning the years between 1967 and 1990. Individual-level data combine several administrative registers covering all adult individuals in Norway from the same period. We use the individual-level data to construct time-series of wage and labor supply by skill groups for each municipality and time period (we consider each municipality to be a different local labor market). We also have information on R&D activities for a subsample of firms between 1970 and 1985, but not for every year in that interval.

The labor market impacts of college openings are established using a synthetic control method (Abadie, Diamond, and Hainmueller 2010). There are many fewer municipalities experiencing a college opening over the period we study than not experiencing a college opening, and this method enables us to find appropriate control municipalities for each municipality with a college opening. Our main results are robust to using a standard difference-in-difference estimator. We model the demand for skilled workers using the data generated from the synthetic control estimator, allowing workers in different age groups to be imperfect substitutes (as in Card and Lemieux 2001).

evidence on the technology channel, and we discuss the human capital spillover literature at end of this section.

Using this model, we quantify the extent to which a college opening induces SBTC.² We note that the opening of a college simultaneously affects the supply of skills (through the production of new graduates), and the demand for skilled workers (indirectly, through endogenous SBTC). To unbundle the influence of these two forces on skill prices, we need an additional assumption. We assume that the impacts of SBTC on the labor market do not take place immediately after the reform, so that only pure supply effects are observed in that period. In our benchmark model, we assume that these impacts do not occur until (at least) two years after the reform. This assumption can be justified if, for example, firms do not invest immediately in response to a college opening but wait until some of the increase in skill supply materializes. It could also be justified if there are delays in the implementation of a new technology.

Our findings are robust to changes in how long the delayed response to SBTC is assumed to be. In addition, we also generate similar results from a model where we do not estimate the elasticity of substitution between skilled and unskilled workers and instead use reasonable estimates of this parameter from the literature. This allows us to avoid having to assume any delay in the SBTC response when estimating the impact of SBTC on the demand for skilled workers.

The impact of the reform on firm-level productivity is estimated using standard production function techniques, and the impact on firms' R&D activities is studied using a typical difference-in-difference estimator. This is because our estimation of the structural parameters of firm-level production functions is not as amenable to a synthetic control estimator as the estimation of reduced-form labor market impacts of the reform. In turn, our R&D data do not allow us to use a sufficient number of pre-reform years needed for a credible implementation of the synthetic control estimator.

We contribute primarily to the literature examining the rising trend in the college premium (e.g. Katz and Murphy 1992; Berman, Bound, and Machin 1998; Machin and Van Reenen 1998; Card and Lemieux 2001; Autor, Katz, and Kearney 2008), and the literature on whether the simultaneous increase in the supply of skilled workers and their wages could be due to endogenous SBTC (e.g. Acemoglu 1998; Beaudry and Green 2003; Blundell, Green, and Jin 2022). We provide new evidence that endogenous SBTC responses to shocks in the supply of skill lead to quantitatively large increases in the skill premium, consistent with a strong relative equilibrium bias (Acemoglu 2007). In the long-run, our results are also consistent with a strong absolute equilibrium bias where the marginal products of skilled workers increase and the demand curves for skilled labor become upward-sloping. Our work is complementary to recent empirical papers on the reaction of directed technical change to changes in factor supplies (Acemoglu and Finkelstein 2008; Hanlon 2015; Bloom, Draca, and Van Reenen 2016). Our focus is, however, on the labor market and our evidence points to endogenous technical change for a large group of workers, which is more general

2. We also explore the implications of other models of local labor markets, including a simple multi-sectoral trade model. We find them to be inconsistent with our empirical findings.

than the existing studies (which tend to focus on specific industries). Our work also complements estimates of technology skill complementarity (Akerman, Gaarder, and Mogstad 2015). If technology and skill are complementary in production, then the increase in the abundance of skills induces firms to invest in technology.

This paper is also related to the literature analyzing local economy adjustments to labor supply shocks. That literature often uses immigration flows in a local labor market as a source of change in factor supplies. Following a positive shock to the supply of skilled labor, the possible types of adjustments are through changes in factor prices (by decreasing skilled wages), changes in product mix (by producing a more skill-intensive product mix), and changes in technology (by adopting or spending more to develop skill-biased technologies). The first channel is by no means unimportant, but in light of some evidence that low-skill immigration has little effect on wages, the recent literature has increasingly focused on the latter two channels of adjustments. For example, a number of papers find that most of the adjustment happens through within-industry changes, which they interpret as changes in production technology (Hanson and Slaughter 2002; Albrecht, van den Berg, and Vroman 2009; Lewis 2011; Peri 2012; Dustmann and Glitz 2015).³ As in these papers, we document that firms adjust their investment in new technologies when faced with a shock to the supply of skilled labor. What is new in our paper is that we document the dynamic impacts of this endogenous technological response on the demand for skilled labor and, consequently, on the wages of skilled workers.

Our paper also closely relates to a growing literature estimating the spillover effects of education.⁴ Most of the earlier papers in this literature attempted to estimate the size of spillovers from education by comparing the wages of otherwise similar individuals who work in cities or states with different average levels of education (e.g. Rauch 1993; Acemoglu and Angrist 2001; Moretti 2004b; Ciccone and Peri 2006). One exception is Moretti (2004a), who estimates education spillover effects on the productivity of manufacturing plants in the US. He finds that the productivity of plants in cities that experience large increases in the share of college graduates rises more than the productivity of similar plants in cities that experience small increases in the share of college graduates. By combining the experimental variation from the college expansion reform and the production function estimation, we extend Moretti (2004a) to estimate the effects of education spillovers on both factor-neutral *and* factor-biased

3. A related strand of literature studies the impact of high-skilled immigration on innovation and productivity of US firms (see Kerr (2013) for a review). Using state-decadal variation in high-skilled immigration, Hunt and Gauthier-Loiselle (2010) find large increases in innovation following upon immigration. Kerr and Lincoln (2010) find an increasing employment of skilled workers in US firms that experience growth of skilled immigrants with H-1B visas. Peri, Shih, and Sparber (2015) further find city-level productivity increases following from H-1B program expansions in local areas that extensively rely on the program. Moser, Voena, and Waldinger (2014) find that Jewish scientist expellees from Nazi Germany to the US encouraged innovation by attracting new researchers to their fields.

4. In the macro-growth literature, recent papers have provided empirical evidence pointing to the large role of human capital externalities in explaining regional variation in development (e.g. Acemoglu and Dell 2010).

productivity. The nature of our panel also allows us to examine the dynamic changes in productivity both in the short- and the long-run.⁵ The spillover effects of education in our paper work through endogenous technology investment responses by firms, which is a possible (but not the most standard) interpretation of education spillovers.

More recently, Kantor and Whalley (2014, 2019) study the spillover effects of research innovations from higher education sector. Kantor and Whalley (2014) find that an exogenous increase in a university's spending increases the local average wage in the non-education sector, particularly among industries that employ more college graduates. Kantor and Whalley (2019) find that proximity to the newly established agricultural experiment stations at land-grant US colleges in the late 19th century had large positive effects on land productivity for two decades. They show that knowledge spillovers increase with literacy rates, consistent with the hypothesis that the knowledge spillovers are skill-biased.⁶ Our paper complement these papers by providing new evidence on how an exogenous increase in supply of skills can affect technology change and local labor market. Our results are driven by an increase in the supply of skilled workers rather than any research innovations produced by the new colleges (which was the focus in the two aforementioned papers). One limitation in our study is that we do not observe the timing of adoption of a specific new technology at the firm level. Although all the evidence from our worker- and firm-level data and R&D data point to endogenous technical change as the most plausible interpretation, we cannot completely rule out the alternative interpretations of human capital spillovers.

This paper is organized as follows. Section 2 provides background on the college expansion reform and a description of the data and sample selection procedures. Section 3 presents evidence on the supply of skills and wages at the local labor market level and interprets the results using the model of Card and Lemieux (2001). In Section 4, we provide plant-level evidence on endogenous technical change via estimating production functions. Section 5 presents further evidence that college openings induced firms to invest more in R&D activities. The last section concludes.

2. Institutional Background and Data

2.1. The College Expansion Reform

The goal of the Parliament when establishing regional colleges in Norway from 1969 onward was to alleviate the increasing problem of capacity at the existing universities.

5. In Moretti (2004a), the panel of plants consists of observation in two years (in 1982 and 1992).

6. Davis and Dingel (2019) provide a theoretical model that could generate human capital externalities that are "skill-biased". In their model, individuals can choose to produce either tradable or non-tradable (service) goods. Crucial to the model are the assumptions where (i) tradable productivity is positively correlated with individual's ability and participation in idea exchanges and (ii) individual ability and local learning opportunities are complements in the production function for tradable goods. These assumptions imply that the output gain from greater ability is increasing in local learning opportunities, generating human capital externalities that are "skill-biased".

There was an increasing demand for college education due to a combination of factors, potentially including population growth, changes in industry composition, and increased mandatory education, implemented from the late 1950s (e.g. Black, Devereux, and Salvanes 2005; Aakvik, Salvanes, and Vaage 2010).⁷

In the mid-1960s, there was strong agreement in the Norwegian Parliament that there was a “national need” to expand the supply of higher education, but the country did not have sufficient resources to build new universities. In 1966, a committee appointed by the government (the Ottosen Committee) proposed to expand the higher education sector by opening regional colleges aiming to provide shorter (two and three years) college education programs. In its report, the Ottosen committee proposed to divide the country into twelve educational regions. Four of the regions already had universities, and the committee proposed that the remaining eight regions should each have one regional college (Ottosen-committee 1966–1970).⁸

Following proposals from the Ottosen Committee, in 1968, the Parliament voted for the opening of four new regional colleges.⁹ For the first batch of new colleges, the Parliament initially agreed to an experimental period of five years, followed by an evaluation. However, in 1970, the Parliament decided to expand the reform to two more regions,¹⁰ and, through the 1970s, the establishment of regional colleges was expanded to all the educational regions.

The report from the Ottosen Committee also suggested three criteria governing the geographic location of a new college within each educational region. First, new colleges should be geographically dispersed across the country. Online Appendix Figure A.2 shows the geographic location of the new colleges across the country.¹¹ This criterion is clearly met for all the new colleges. Second, new colleges had to be established in regions and in municipalities where all the necessary infrastructure could be completed

7. Online Appendix Figure A.1 shows that the overall educational level increased dramatically in Norway from 1960 to 1990. In 1960, only 4.2% of the population above the age of 16 had a college or university degree and 16.3% had a high-school diploma, so the ratio of college-educated people to high-school-educated people was about 0.26. During the decade we study, the relative supply of college-educated workers rose from 0.28 in 1970 to 0.35 in 1980. This increase in the relative skill ratio was much larger than what was observed either in the preceding decade (from 1960 to 1970, when this ratio increased from 0.26 to 0.28), or in the following decade (from 1980 to 1990, when this ratio increased from 0.35 to 0.38).

8. In Norway, reports from expert committees are followed up by white papers from the government with explicit suggestions for the Parliament to vote on. In the case of the higher education sector over the last several decades, there has been one of these reports from expert committees about every decade. This is the main procedure determining the total amount of resources allocated to the higher education sector and resource allocations to each college.

9. Three of them were opened starting in the fall of 1969, located in the counties of Rogaland (in the municipality of Stavanger), Agder (in the municipality of Kristiansand), and Møre og Romsdal (dual locations in the municipalities of Volda and Molde). The fourth college was opened in the fall of 1970 in the county Telemark (located in the rural center of Bø).

10. These are located in the educational region of Hedmark/Oppland (in the municipality of Lillehammer), and one in the region of Nordland (in the municipality of Bodø).

11. Note that, in this figure, and in our subsequent analysis, we also include three colleges that were built in the same period but were not part of the recommendation of the Ottosen Committee. See the Data Section for details.

within a year of the establishment decision. Since this was a necessary requirement to get started within the given time limit, it is clear that this criterion was also met. Third, new colleges should be placed to stimulate growth in regions with stagnation problems. We do not find much empirical support for this last criteria (see Table 1 where we show the mean of sector compositions, manufacturing output, and labor market outcomes in the non-treated areas, together with the difference between the treated areas and non-treated areas). Our interpretation is that the two first criteria dominated the selection for placement of the colleges.

In Online Appendix Section A, we provide additional details on the characteristics of the new colleges in terms of their sizes and academic programs. We also show that there was very little research output produced by these regional colleges in the period under study.

2.2. Data and Sample Selection

We use both firm- and worker-level data from several sources covering the period 1967–1990. Below, we describe the data we use.¹²

Worker-Level Data. Our worker-level data come from two sources. The first one contains the data on workers from administrative registers prepared by Statistics Norway. The data cover all Norwegian residents aged 16–74 years old, covering the same years as the plant-level data (1967–1990). The variables captured in this dataset include individual demographic information (such as sex and age) and socioeconomic data (such as completed level of schooling, municipality of residence, and annual earnings). For certain male cohorts, we have data on their IQ scores upon entering military service.¹³ In addition to the administrative registers, we also use the Norwegian Census, which was conducted in 1960, 1970, and 1980. The census covers the entire population and has additional information on labor market activities (such as industry of employment). A unique personal identifier allows us to follow workers over time and to link the census data with the registry data.

Our wage measure is based on men's annual labor earnings from the administrative registers.¹⁴ Annual labor earnings are the sum of pre-tax labor income (from wages and self-employment) and work-related cash transfers (such as unemployment benefits and short-term sickness benefits). For the period we study, it is not possible to separate

12. The earliest year of our plant- and worker-level data begins in 1967. It could be potentially interesting to expand the data to analyze the effects of the reform in more recent years. However, following the appointment of a new committee (the Hernes Committee), a new round of reforms was initiated in the early 1990s. By the mid-1990s, regional colleges were consolidated and upgraded to university colleges, where they were given the right to develop research-based degrees, hire professors, take part in the training of researchers, and engage in fundamental as well as applied research. For this reason, we limit our sample period to 1990. During this period, there were no major reforms to the higher education sector.

13. We explain the use of these data in Online Appendix D.

14. Similar to Card and Lemieux (2001), we focus on male wages because there is a large increase in female labor supply over the period under study.

TABLE 1. Comparison of baseline characteristics before the reform by treatment status.

Outcomes	Worker-level data			Firm-level data	
	Non-treated	Difference	Outcomes	Non-treated	Difference
Share of skilled workers, young	0.129*** (0.002)	0.002 (0.007)	Log output per worker	5.436*** (0.010)	-0.017 (0.031)
Log high-skilled wages, young	12.387*** (0.006)	0.026 (0.018)	Plant size	48.202*** (1.670)	-1.602 (5.203)
Log low-skilled wages, young	12.169*** (0.005)	-0.000 (0.014)	Hours worked	73195*** (2499.6)	-1752 (7787.0)
Share of skilled workers, old	0.099*** (0.003)	0.021* (0.009)	Log total wage costs	14.939*** (0.020)	0.018 (0.061)
Log high-skilled wages, old	12.794*** (0.004)	0.012 (0.013)	Investment in machinery	809.656*** (31.803)	-115.901 (99.067)
Log low-skilled wages, old	12.271*** (0.007)	0.012 (0.021)	Investment in machinery (excluding transport)	667.045*** (28.932)	-68.377 (90.125)
Growth in skill shares, young	0.002*** (0.001)	0.002 (0.002)	Investment in machinery and facilities	1146.9*** (47.869)	-21.370 (149.116)
Growth in high-skilled wages, young	0.002 (0.004)	0.007 (0.012)	Employment by high-skilled industries (%)	0.483*** (0.015)	0.047 (0.045)
Growth in low-skilled wages, young	-0.016*** (0.001)	0.001 (0.003)	Plants in high-skilled industries (%)	0.374*** (0.010)	0.016 (0.031)
Growth in skill shares, old	0.002*** (0.000)	0.001 (0.001)	Outputs by high-skilled industries (%)	0.493*** (0.015)	0.051 (0.047)
Growth in high-skilled wages, old	-0.015*** (0.002)	0.002 (0.006)	Log total costs of R and D	7.814*** (0.260)	-1.554** (0.718)
Growth in low-skilled wages, old	-0.014*** (0.001)	-0.001 (0.003)	Log total R and D man-years	3.336*** (0.256)	-1.621** (0.707)

Note: This table reports the mean pretreatment characteristics of the municipalities with treatment and the remaining municipalities without treatment. For growth rate outcomes, the table reports the annual change between 1967 and 1968. For the rest of the outcomes, the table reports the averages in 1967. All means are weighted by the number of plants in 1967. * p < 0.10, ** p < 0.05, and *** p < 0.01.

the two. The Norwegian earnings data have several advantages over those available in most other countries. First, there is no attrition from the original sample because earnings data come from tax records and tax records are in the public domain in Norway. Second, our earnings data pertain to all individuals, and not only to jobs covered by social security. Top-coding is only performed at very high earnings levels, which is another advantage over other similar data sets such as the social security data in the US. However, the top-coded amount increased in the 1970s and early 1980s. To make sure that our average earnings series are comparable over time, we manually top-code the earnings data for all earnings above the 96th percentile of the earnings distribution by year.¹⁵

Education attainment is reported by the education authorities directly to Statistics Norway, thereby minimizing any measurement error due to misreporting. For every individual, the data record the year of graduation for each level of completed education. Based on this information, we measure the highest level of completed education for each individual in each year.¹⁶

The sample of individuals being analyzed includes workers aged 20–62 years and whose annual earnings are at least twice as large as the basic amount required to participate in the national social insurance program.¹⁷ In each year, we classify workers into two skill groups. The high-skilled group includes workers who have completed at least some college education. The unskilled group consists of all remaining workers. Our definition of a local labor market is a municipality, which is the smallest administrative unit in Norway.¹⁸ By combining workers' information on skill levels and municipality of residence, we construct measures of skill composition and wages in a local labor market over time. For instance, the share of high-skilled

15. Top-coding exists in all years prior to 1986, except for 1981. The 96th percentile of earnings distribution by year is not subject to top-coding across all years, which we use as our top-coding thresholds. For earnings above the top-coding threshold, the top-coded amount we use is equal to the earnings at the 96th percentile in each year times 1.32, where 1.32 is obtained by dividing the average earnings above the top-coding threshold in 1981 by the value of that threshold (there is no top-coding in 1981). In the previous version of the paper, we only trim the top 0.1% observations with the highest earnings by year. The top-coding amount (i.e. the censoring point) went up starting in 1971, which led to a discrete jump in the mean earnings, and the jump is bigger for skilled-earnings than unskilled-earnings in that year (because a higher fraction of skilled workers were subject to top-coding prior to 1971). Because year 1971 corresponds to the second year post reform for five new colleges, the previous version of the paper documented a discrete jump in mean and relative earnings in both the treated and synthetic groups in year 2, which is in fact driven by changes in top-coding in the data.

16. The educational establishment data are available starting from 1970. Information on the year of graduation is also left-censored in the year 1970. The completed education levels in the years 1967–1969 are imputed using the completed education level in 1970.

17. Although the mandatory retirement age is 67 years, about 80% of Norwegian workers are entitled to receive early retirement benefits beginning at age 62 years (Bhuller, Mogstad, and Salvanes 2017). The basic amount adjusts for costs of living each year. We validated our sample of workers by linking our earnings data with the 1970 census, which contains categorical information on annual hours of work. More than 96% of individuals in our data are full-year workers.

18. Our definition of the local labor market is consistent with previous empirical work that relies on geographical segmentation of the Norwegian labor market (e.g. Akerman, Gaarder, and Mogstad 2015).

workers in year t and municipality c is given by the ratio of the number of high-skilled workers over the total size of the labor force residing in municipality c and year t , and the mean wage among high-skilled workers is defined as the average log annual earnings among skilled men residing in municipality c and year t .

Plant-Level Data. Our main plant-level data are drawn from the Manufacturing Statistics collected annually by Statistics Norway for the period of 1967–1990. The Manufacturing Statistics covers *all* plants in the mining, quarrying, and manufacturing sectors operating during the calendar year in Norway.¹⁹ The response rate is extremely high because firms are required by law to submit their survey responses.²⁰ A consistent and unique ID on each establishment allows us to create a panel of plants over this period. We focus on plants in the manufacturing sector with more than five employees, completing at least a total of 5,000 hours worked in a year. The restriction on size is driven by the fact that complete questionnaires were only collected from plants having at least five employees.²¹ The restriction on total hours ensures that the plants in our sample are active in production in any given year.

The firm-level data contains detailed information on output, inputs, and production costs. Using this information, we compute value-added per firm and year, defined as the gross value of production minus the costs of materials and services.²²

To measure capital stock, we use the fire insurance value of buildings and equipment owned by the firm and yearly investment flows. The fire insurance value of capital stock is available only from 1974 on. Furthermore, the nature of the fire insurance value means that there is not much variation over time in this variable. Therefore, for each plant, we take the fire insurance value as the value of capital in the first year of the panel and impute the value of the capital stock in subsequent years by adding up the value of net investment (in buildings and equipment) in each year and assuming that the current equipment and buildings depreciate at a constant rate.²³ For plants for which we do not observe the fire insurance value in the first year of the panel (plants which first appeared before 1974 when the fire insurance value was not available), we take the mean fire insurance value by municipality–industry cells in 1974 (separately

19. An establishment is defined as a functional unit that, at a single physical location, is engaged mainly in activities within a specific activity group. A firm doing business in different municipalities is shown as two or more separate establishments in the sample. In the Norwegian context, most plants belong to separate firms (Klette and Griliches 1996). See Halvorsen, Jensen, and Foyn (1991) for a detailed description.

20. The questionnaires were usually sent out in April/May after the end of the reference year, with a response deadline of four to five weeks. Firms failing to respond to the initial inquiry were sent written follow-up letters for up to six months from the first deadline. Firms that did not respond by then were fined.

21. For small plants with less than five employees, information was extracted from separate administrative registers, which contained fewer variables than the original questionnaire.

22. Value-added is measured at factor prices, defined as the gross value of production (value of gross output, including subsidies), less the cost of goods and services consumed (excluding VAT) and indirect taxes (except VAT and investment levy).

23. The discount rates being used are 0.05 for equipment capital and 0.02 for buildings.

by buildings and equipment), and use the corresponding cell-specific means as the initial capital stock in the first sampling year of the firm.

To measure labor input, we use the total hours of employment for each plant. Unfortunately, from the firm-level data, we cannot distinguish labor inputs by skill groups. In addition, for the time period under study, we are not able to link the worker-level data with the plant-level data (this only becomes possible after the 1980s). For this reason, our analysis of worker productivity is conducted at the industry and municipality level, since we can observe the skill composition of the labor force at these more aggregated levels by combining worker- and plant-level data. This is explained in detail in Section 4.

Firm-Level R&D Data. We also have information on R&D activities for a subsample of firms. During this period, information about R&D is collected from R&D surveys conducted by the Royal Norwegian Council for Scientific and Industrial Research. The R&D sample includes mainly manufacturing firms above a certain size class.²⁴ We have access to data starting from 1970 and then biannually from 1975 to 1985. The R&D data can be linked to the plant-level data using a combination of firm and detailed industry identifiers. In Section 5, we discuss our use of the R&D data in detail.

College Reform Data. The main source of college reform data is from Ottosen-committee (1966–1970), annual National Budgets (with details on financial support for each college, including the number of students), and Johnsen (1999), which contains detailed information on the timing, location, programs, and student enrollment of all new regional colleges. Twelve new colleges were built out of the reform initiative in the period we study. We also carefully checked against opening dates of all colleges in Norway and included three additional colleges that were built in the same period but were not part of the recommendation from the Ottosen Committee (and are similar to the colleges originally mentioned in the report). The first college opening occurred in year 1967 and the last college opened in 1981.

3. Worker-Level Evidence: Wages, Skill, and Skill-Biased Technical Change

We begin this section by documenting positive impacts of a college opening both on the supply of skills and relative wages of skilled workers in local labor markets. We interpret our estimates using the model of Card and Lemieux (2001). Our results suggest that for several years following the establishment of a college, there is significant skill-biased technology change in the affected labor markets. In Section 4, we provide more direct evidence of endogenous technical change by quantifying the effects of the

24. The size limits varied among different sectors. The size limits were lower in sectors known to be R&D intensive (down to 10 employees) and higher in sectors with low R&D activity (up to 100 employees). For instance, in the machinery and equipment industries utilized, the R&D surveys have close to full coverage for firms with more than 20 employees. For detailed description of the data, see Møen (2005).

reform on labor productivity, by estimating production functions on firm-level data, as well as estimating the impact of the reform on firms' R&D investments.

3.1. Construction of the Control Group

There are only 15 municipalities that benefited from a college opening during the reform period we consider, out of a total of nearly 400 municipalities. In principle, all untreated municipalities can be potential control municipalities, but the danger of proceeding this way is that only a few of them may be similar to the relatively small set of reform municipalities in the treatment group.²⁵

Therefore, we select comparison municipalities for the control group using the synthetic control estimator developed in Abadie, Diamond, and Hainmueller (2010) (hereafter **ADH**).²⁶ For each municipality with a college opening, we use the **ADH** method to construct an optimal synthetic control group, with pretreatment trends in the outcome variable matching those of the treated municipality. This method is suitable in our setting where a discrete treatment (i.e. a new college) is applied to one unit (i.e. a municipality) and not to others, within a large geographic area. The idea is to select control groups based on a set of pre-intervention characteristics Z_{it} , which predict the outcome of interest after the treatment, where Z_{it} includes pre-reform (time-varying) outcome variables (such as the whole history of outcomes), as well as pre-reform (time-invariant) characteristics of the municipality. This procedure provides a vector of municipality-specific weights that minimize the distance between the treated municipalities Z_{it} and the weighted mean of the synthetic control.

In our setting, Z_{it} includes the outcomes measured in each of the five years prior to the treatment, normalized by the outcome in the year of the treatment.²⁷ Z_{it} also includes a set of municipality-level characteristics averaged over pre-reform years, including demographic composition (share of workers aged 20–35 years among the workforce), and skill composition of the labor force (share of high-skilled workers). As a result, for each outcome, the pre-reform trend (the change in the outcome variable in each of the five years prior to the reform), and the skill and age compositions of the labor force in the synthetic control municipality should track closely those in the treated municipality. Because Z_{it} contains pretreatment outcome variables, a different synthetic control is used for each outcome. To make sure that the control municipality

25. Table 1 shows the characteristics between the municipalities with a new college and the remaining municipalities prior to the reform. The log average wages by skill groups are fairly close between treated and remaining municipalities. Relative to the non-treated municipalities, it appears that treated municipalities comprised of a more educated labor force and also experienced faster growth in the skill shares.

26. As we discuss below, most of our results are robust to using instead a standard difference-in-difference estimator.

27. In cases where the pre-reform period is less than five years, Z_{it} includes the pretreatment trends in the outcome variable in all available years prior to the reform.




FIGURE 1. The effects of the reform on relative wages and skill compositions. This figure presents the synthetic control estimates on skill composition and relative wages of the workforce. On each graph, the year of the reform is normalized to period zero. Each data point represents the mean outcome in a given period, relative to the levels in the year of the reform (where the outcome in the reform year is normalized to zero). Each graph reports the weighted average of all treated municipalities and the corresponding synthetic controls, with weight given by the number of plants in the treated municipality.

is geographically similar to the treated municipality, we restrict potential control municipalities (“donor pool”) to be in the same region as the treatment municipality.²⁸

3.2. Effects of the Reform on Skill Compositions and Wages

3.2.1. Main Results. Figure 1 presents the effects of the reform on the skill composition of the workforce, and the relative earnings of skilled versus unskilled

28. We divide the municipalities into four geographical regions as follows: North (Finnmark, Troms, Nordland), Middle (Nord-, Sør-Trøndelag, Møre og Romsdal), West and South (Sogn og Fjordane, Hordaland, Rogaland, Aust- og Vestagder), and East (Telemark, Vestfold, Buskerud, Oppland, Hedmark, Oslo, Akershus, Østfold). Although we think the municipality level is a reasonable approximation of local labor markets and ensures comparability with previous empirical work that relies on geographical segmentation of the Norwegian labor market (see e.g. Black, Devereux, and Salvanes 2005; Akerman, Gaarder, and Mogstad 2015), we check how spillover effects can affect our results. To this end, we exclude the municipalities within 30 km of the treated municipality from the “donor pool”, and re-estimate how college openings affect relative wages and supply using the synthetic control method. We find that the results do not materially change, demonstrating the estimated effects of the reform are not driven by comparison between the treated municipalities and nearby municipalities where there is the most commuting (the full results are available upon request).

workers, for treatment and control local labor markets (analogous estimates for absolute levels of earnings of skilled and unskilled workers are presented in Figure A.5 in the Online Appendix). Because workers in different age groups are possibly imperfect substitutes in production (Card and Lemieux 2001), we split the sample into young (aged below 35 years) and old (aged above 35 years) workers, and analyze the impacts of the reform separately for each group. Workers in the older group may be relatively shielded from the supply effect because the inflow of newly college-educated workers is driven almost exclusively by the young. In turn, workers in the younger age group are affected by both supply and technological effects of the reform.

In each panel of Figure 1, the year of the reform for each municipality is normalized to period zero. For every treatment or control municipality, we compute the difference in the outcome of interest in a given year relative to the level of that variable in the year of the reform (the level of the outcome in the reform year is also normalized to zero). Each panel in Figure 1 then shows the weighted average of these differences across all 15 sets of treated municipalities (thick line) and the corresponding synthetic controls (dashed line), with weights given by the number of plants in the treated municipality in every year. The effect of the reform in each year (after year zero) is the difference between the two lines in each panel. Details of our implementation of the ADH procedure are described in Online Appendix B.²⁹

The top-left panel of this figure shows that, compared with the synthetic controls, labor markets with a new college experience an increase in the supply of skilled young workers. The gap between the treated group and the synthetic control increases over time (because an additional flow of new graduates is added to the stock of skilled workers each year), reaching nearly four percentage points 10 years after the opening of the college.³⁰ In contrast, the reform has little impact on the skill composition among workers aged 35 years or more for the first 10 years following the reform. The share of skilled workers among older workers begins to increase toward the end of the panel, partly because of aging of the cohorts affected by the reform.

29. The treatment effects reported in the synthetic control figures, such as Figure 1, are constructed by averaging municipality-specific treatment effects, one for each of the treated municipalities. The composition of the treated municipalities varies by pre-reform years, which may impact our results. For instance, in Figure 1, the relative wages and skill shares in the treated municipality and the synthetic control in period -5 are computed only for municipalities, which had a new college opened from year the 1972 and onwards, whereas in period -2, all municipalities that had a new college are included. As a robustness check, we also re-estimate the effects of the reform on wages and skill composition using the synthetic control method, where we match the pre-reform outcomes using data only up to 2 years prior to the reform (for all treated) municipalities, where the composition of the treated municipalities is fixed in each of the pre-reform years (see Online Appendix Figure A.6). We find effects on relative wages and skill shares that are very similar to what we report in Figure 1.

30. When we further decompose the unskilled into workers with at least some high school and workers with less than high school, we find that the share of workers with some high school gradually decreases over time in treated markets. This is, in some sense, expected given that people who are on the margin of going to college are more likely to be affected by the opening up of new colleges and, hence, move from the middle-skilled category to the high-skilled category.

The lower panel of Figure 1 shows estimates of the effects of the reform on the relative earnings of skilled workers by age groups. Immediately after the college opens, the relative earnings of young skilled workers are slightly higher in the control group than in the treatment group. Toward the end of the period we study, the relative earnings of skilled workers in the treatment group increase above those in the control group. Among older workers, the relative earnings of skilled individuals increase substantially following a college opening (this pattern is also seen for absolute earnings of skilled workers, as shown in Figure A.5 in the Online Appendix).³¹

A reasonable interpretation of these findings is the following: The drop in the relative wages of young skilled workers immediately after the reform is a consequence of increased supply putting downward pressure on wages. However, the fact that, in spite of the rise in supply, the relative wages of skilled young workers in treatment areas eventually rise above those in control areas, suggests that the demand for these workers increased more rapidly in treatment than in control areas. This could be possibly caused by an endogenous increase in investment in skill-biased technologies by firms in treatment areas since the abundance of skilled workers may have made the use of these technologies more profitable.³² This would be consistent with our findings that older skilled workers experience a stronger increase in their relative wages in treatment areas than in control areas. Older skilled workers are shielded from any downward pressure on their earnings due to the increase in the supply of skills, because young and old workers are imperfect substitutes. Their relative earnings increased because the reform increased the demand for skilled labor among older workers without affecting the supply. In Section 3.3, we develop in more detail this interpretation of our findings.

To assess the extent to which our estimates are statistically important, we follow Abadie, Diamond, and Hainmueller (2010) and estimate a series of placebos by iteratively applying the synthetic control method to every municipality in the pool of potential control municipalities. In each iteration, we reassigned a treatment from a treated municipality to a control municipality (for details, see Online Appendix B). This procedure is repeated for each treated municipality so that, for each of them, we obtain

31. Norway is characterized by a centralized wage bargaining system where wages are partially set at the national level through the central employer's and employee's organizations. This implies that there is a stronger degree of wage compression than in most other countries (Moene and Wallerstein 1997; Kahn 1998). Moreover, in the 1970s, the Norwegian government more actively took part as a third party in the central negotiations with the aim of reducing the nominal increases in wage rises, resulting in very little (if any) real wage growth (Aanensen 2010). Importantly, though, since we do not rely on the national variation as a source of identification, these national changes are in principle differenced out in our estimates.

32. One potential challenge to the interpretation of the wage effects of young workers is that changes in the relative number of college graduates might affect the relative composition of the pool of college graduates. For instance, after the reform, selection into college- education may be based on ability to a greater extent than mobility costs. To address this concern, we use the IQ information of several cohorts of males from the military draft data. We do not find any evidence that the reform changes the average cognitive ability (proxied by IQ) among college and non-college workers. See Online Appendix D for details.




FIGURE 2. The effects of the reform on relative wages and skill compositions: placebo tests. This figure shows the results for the placebo test. The gray lines represent the estimated treatment effect from each placebo of the permutation test. The thick solid line denotes the treatment effect estimated using the actual treated municipalities in the data. The implied p -values of the actual treatment effects are reported in Online Appendix Table A.1. The outcome in the year of the reform is normalized to zero.

an empirical distribution of the estimated gaps between the “treated” municipality and its synthetic control.

In principle, we can calculate municipality-specific p -values for the test that the treatment effect is zero. However, it is simpler to present a single p -value for the treatment effects averaged across all treated municipalities. We begin by randomly drawing 50 placebos (with replacement) for each treated municipality and average them across all treated municipalities. We then calculate p -values based on the distribution of the treatment effects from these aggregated placebos.

Figure 2 shows the results of this procedure. The gray lines represent the year-by-year treatment effects for each placebo. The solid black line denotes the treatment effect estimated using the actual data (from Figure 1), with the observed treatment assignment. The implied p -values for each of the actual gaps in each year, that is,

the proportion of placebo gaps larger than the estimated gap, are presented in Online Appendix Table A.1.

There are two variables for which the estimated treatment effects are large and statistically important almost every year after the reform: skill shares among young workers and the relative earnings of skilled older workers. When we consider outcomes many years after the reform, there are statistically significant treatment effects for all outcomes considered in Figures 1 and 2 (Figure A.7 in the Online Appendix presents the permutation tests for earning levels and Online Appendix Table A.1 shows the corresponding *p*-values).

Instead of the ADH procedure, we could have used a standard difference-in-differences research design. In Online Appendix C, we report findings from this exercise where all the untreated municipalities are included as comparisons. The identifying assumption underlying the regression analysis is that the geographic location of the college expansion is not correlated with different underlying trends in local labor-market outcomes across the markets (common trends). As a first check of whether this is a plausible assumption, we examine whether the outcome variables in the treated and control regions have similar trends over time during the pre-reform period. For certain outcomes, the pre-reform trends appear to be different between the treated regions and the remaining areas used as comparison. Therefore, the synthetic control group may be especially helpful in our case for identifying which municipalities should go into the control group. Nevertheless, the effects of the reform on aggregate skills and relative wages across the two age groups are qualitatively similar to the synthetic control estimates.

3.2.2. STEM versus Non-STEM Colleges. Out of the 15 new colleges in our study, 9 had majors in STEM fields (we call these STEM colleges). In this section, we ask whether the impacts we estimate are due mainly to openings of STEM colleges because it is plausible that STEM graduates are the ones whose productivity most responds to technical change, and conversely, STEM graduates may provide the greatest incentive to firms to upgrade their technology. Of course, the decision of whether to offer any STEM majors is endogenous and may depend on the existing (pre-reform) local industrial structure, so our estimates have to be interpreted with caution.

Figures 3 and 4 present the estimated effects of reform for STEM and non-STEM colleges for young and old workers, respectively. We find that the opening of both types of colleges led to an increase in the share of young skilled labor in the local labor market, with a stronger effect for STEM colleges. However, the relative earnings responses reported in Figure 1 are driven exclusively by those regions where a STEM college was established.³³

33. When we aggregate the synthetic control estimates from each individual reform municipality, we use the number of plants in the municipality as a weight. For municipalities with a non-STEM college, the average number of plants is smaller than the municipalities with a STEM college. Therefore, STEM




FIGURE 3. The effects of the reform on young workers: STEM versus non-STEM colleges. This figure presents the synthetic control estimates on skill composition and wages of young workers, separately for STEM colleges and non-STEM colleges. On each graph, the year of the reform is normalized to period zero. Each data point represents the mean outcome variable in a given period, relative to the levels in the year of the reform (where the outcome in the reform year is normalized to zero). Each graph reports the weighted average of all treated municipalities and the corresponding synthetic controls, with weight given by the number of plants in the treated municipality in the given year.

Labor markets where there was an opening of a non-STEM college did not experience substantial wage changes relative to control labor markets. The relative wages of skilled old workers move similarly in control and treatment areas. Similarly, in the years immediately following the opening of a college, the relative wages of skilled young workers do not follow distinct paths in treatment and control areas, perhaps because the inflow of skilled workers is not large enough to make a difference at the very beginning, or because it takes some time for wages to adjust. However, it is remarkable that several years after the reform, the relative wages of skilled young workers are much lower in treatment areas than in control areas, probably due to the downward pressure on wages resulting from an increase in the supply of skilled workers in the medium run (and the absence of a technological response). The results

colleges carry a larger weight when it comes to the overall effects, contributing to the similarities between Figures 3 and 1.




FIGURE 4. The effects of the reform on old workers: STEM versus non-STEM colleges. This figure presents the synthetic control estimates on skill composition and wages of old workers, separately for STEM colleges and non-STEM colleges. In each graph, the year of the reform is normalized to period zero. Each data point represents the mean outcome variable in a given period, relative to the levels in the year of the reform (where the outcome in the reform year is normalized to zero). Each graph reports the weighted average of all treated municipalities and the corresponding synthetic controls, with weight given by the number of plants in the treated municipality in the given year.

concerning earnings levels (as opposed to relative earnings), which show an increase in the earnings of skilled workers in areas where there was an opening of a STEM college, are presented in Online Appendix Figures A.8 and A.9.

3.3. Separating Supply and Demand Factors

3.3.1. The Model Setup. We use the model in Card and Lemieux (2001) to decompose the differences in trends in skill- and age-specific wages between reform and non-reform areas (reported in Online Appendix Figure A.5) into supply and technology factors. Assume that aggregate output in period t and labor market D depends on two constant elasticity substitution (CES) sub-aggregates of skilled (college) and unskilled (non-college) labor:

$$Y_{t(D)} = (\alpha_{t(D)}(a_{t(D)}S_{t(D)})^\rho + (1 - \alpha_{t(D)})(b_{t(D)}U_{t(D)})^\rho)^{\frac{1}{\rho}} \quad (1)$$

and

$$S_{t(D)} = \left[\sum_j \beta_j^s S_{jt(D)}^\eta \right]^{\frac{1}{\eta}}$$

$$U_{t(D)} = \left[\sum_j \beta_j^u U_{jt(D)}^\eta \right]^{\frac{1}{\eta}}.$$

The gross elasticity of substitution between different age groups j with the same level of skill is $\sigma_A = 1/(1 - \eta)$, where $\eta \in (-\infty, 1)$. Workers of different ages are gross substitutes when $\sigma_A > 1$ (or $\eta > 0$) and gross complements when $\sigma_A < 1$ (or $\eta < 0$). If different age groups within a given level of skill are perfect substitutes, then η is equal to 1.³⁴ $\sigma_E = 1/(1 - \rho)$, where $\rho \in (-\infty, 1)$, is the elasticity of substitution between skilled and unskilled workers and substitutes, and complements are defined as above. (These elasticities of substitution determine how changes in technology and changes in the supply of skills across cohorts affect the demand for skills and the relative wages of skilled workers.) β_j^s and β_j^u are efficiency units of skilled and unskilled labor of age group j , respectively.

Note that this formulation of the CES production function allows factor-augmenting technologies to affect the productivity of workers through the efficiency units of labor.³⁵ $a_{t(D)}$ and $b_{t(D)}$ represent skilled and unskilled labor augmenting technological change, and $\alpha_{t(D)}$ can be interpreted as indexing the share of work activities allocated to skilled labor (Autor, Katz, and Kearney 2008). Skill-biased technical change involve increases in $a_{t(D)}/b_{t(D)}$ or $\alpha_{t(D)}$.³⁶

Let $D = 1$ denote the treatment group and $D = 0$ denote the synthetic control group. D can affect inputs ($S_{jt(D)}$ and $U_{jt(D)}$) and the technology parameters ($a_{t(D)}$, $b_{t(D)}$, $\alpha_{t(D)}$) in the post-reform periods ($t \geq 1$). The technology response may be an

34. We have assumed that the elasticity of substitution across age groups is the same for both skill groups. In Online Appendix E.1, we show how the model can be extended to allow for separate elasticities across skill groups. Estimation results from the extended model are very close to the model we present in this section.

35. One notable omission from this model is capital. As emphasized, for example, in Beaudry and Green (2003) and Beaudry, Doms, and Lewis (2010), a decrease in the price of skill could induce an endogenous increase in the capital stock if capital and skilled labor are complements. We introduce capital in the empirical model of Section 4, although, in that specification of the production function, technical change is not allowed to impact the elasticity of substitution between capital and skilled labor. If technical change also makes capital and skills more complementary, as suggested in Beaudry and Green (2003), then we may be overstating the impact of endogenous skill-biased technical change on the relative earnings of skilled workers.

36. In the literature, an increase in $a_{t(D)}/b_{t(D)}$ represents intensive SBTC and an increase in $\alpha_{t(D)}$ represents extensive SBTC (Johnson 1997). Intensive SBTC makes skilled labor more productive at the tasks it already performs without replacing the tasks of unskilled labor. Extensive skill-biased technological change can be interpreted as changes in production processes such that skilled workers are profitably employed in some jobs that unskilled workers used to do. This type of technical change is a productivity gain by one factor and a loss by another.

endogenous response to changes in $S_{jt(D)}$ and $U_{jt(D)}$. In period 0 ($t = 0$; pre-reform period), the treated group and the synthetic control have identical labor inputs and productivity parameters. After period 0, labor inputs and technology parameters may differ between the treated group and the synthetic control.

Assuming competitive labor markets (wage equals the marginal product of labor), the ratio of wages for skilled and unskilled workers in each age group j is

$$\begin{aligned} \log \frac{w_{jt(D)}^s}{w_{jt(D)}^u} &= \underbrace{\left[\log \left(\frac{\alpha_{t(D)}}{1 - \alpha_{t(D)}} \right) + \rho \log \left(\frac{a_{t(D)}}{b_{t(D)}} \right) \right]}_{\equiv \log \left(\frac{\theta_{st(D)}}{\theta_{ut(D)}} \right)} \\ &+ \log \frac{\beta_j^s}{\beta_j^u} - \frac{1}{\sigma_E} \log \frac{S_{t(D)}}{U_{t(D)}} - \frac{1}{\sigma_A} \left(\log \frac{S_{jt(D)}}{U_{jt(D)}} - \log \frac{S_{t(D)}}{U_{t(D)}} \right). \quad (2) \end{aligned}$$

Note that the aggregate supply of skills ($S_{t(D)}$) and ($U_{t(D)}$) is unobserved and depends on age-specific skill supplies ($S_{jt(D)}$) and ($U_{jt(D)}$) and the parameters in the sub-aggregate CES production function ($\eta, \beta_j^u, \beta_j^s$).

Let $d \log(S_{1t}/U_{1t}) \equiv \log(S_{1t(1)}/U_{1t(1)}) - \log(S_{1t(0)}/U_{1t(0)})$ denote an increase of the relative supply of young workers in the treatment group relative to the control group. Equation (2) shows that such an increase in the relative supply of young workers can have two direct effects. The first comes from changes in the aggregate relative supply, $-(1/\sigma_E)d \log(S_t/U_t)$. This effect is the same for young and old workers and depends on the elasticity of substitution between skilled and unskilled workers. The second comes from the direct negative effect on relative wages among young workers, given by the term $-(1/\sigma_A)(\log(S_{jt(D)}/U_{jt(D)}) - \log(S_{t(D)}/U_{t(D)}))$. For old workers, a rise in the relative supply of young workers may raise their relative wages (by $(1/\sigma_A)d \log(S_t/U_t)$). In the extreme case where young and old workers are perfect substitutes within skill groups (i.e. $\sigma_A \rightarrow \infty$), then the effect on young workers is identical to the effect on old workers.

The above two effects are standard and due to supply factors. Our model allows for a third effect of the reform on relative wages due to technology change. Our main goal is to estimate the sequence of $d \log(\theta_{st}/\theta_{ut}) = \log(\theta_{st(1)}/\theta_{ut(1)}) - \log(\theta_{st(0)}/\theta_{ut(0)})$, using the data generated in Figure 1.³⁷ These parameters identify the differential technical change in the treated group relative to the control. Given that we assume that the only direct effect of the reform on the economy is a change in the relative supply of skills, the only way technology could have responded in this model is if firms invested in technology upgrading as an endogenous response to changes in skill supplies. Therefore, we say there is evidence of *endogenous* skill-biased technical change if the treated group experienced more accelerated skill-biased technical change

37. We estimate this equation in the generated synthetic control data rather than in the raw data because the synthetic control procedure allows us to construct a good control group for the treatment firms, which we would not be able to replicate by fitting the model directly to the raw data.

than the synthetic control, for the relative supply of skills increased more in the treated group than the synthetic control.

It is, however, possible that the reform could have affected technical change for reasons unrelated to the supply of skills. For example, if new colleges engaged in R&D activities, then they could foster an increase in the amount of innovation being produced at any point in time. In Section 3.3.5, we discuss why this and other alternatives can be ruled out.

3.3.2. Identification and Estimation Strategies. When estimating equation (2), we face two empirical challenges. First, to credibly identify σ_E and σ_A , we need exogenous variation in skill supplies. As argued in the previous section, combining information on college openings with the construction of synthetic controls, we are able to observe arguably exogenous variation in the supply of skills.

Second, any exogenous variation in skill supplies also has an effect on technical change through the channel we emphasize in the paper. Therefore, college openings affect wages through two channels: the direct impact of skill supplies on wages through σ_E and σ_A and the indirect impact of skill supplies on wages through $\theta_{st(D)}/\theta_{ut(D)}$. Using college openings alone as an exogenous shock is not enough to separately identify these two mechanisms. Below, we explore additional assumptions, which allow us to separate these two effects.³⁸

We estimate equation (2) in two steps, using the data generated in Figure 1.³⁹ Below, we discuss the identification arguments in each of the two steps, leaving details of the empirical implementation in Online Appendix E.

In the first step, σ_A (the gross elasticity of substitution between different age groups j within a given skill group) is identified from exogenous shifts in age-specific skill supplies. In particular, we explore exogenous changes in supply (and, therefore, identify the demand curve) by using differences in relative supplies within age (j) and across treatment groups (D) and correlate them to differences in relative wages.⁴⁰ Given the estimate of σ_A , the efficiency parameters β_j^s and β_j^u (which are assumed to be invariant to D) are estimated using the equations derived by equalizing the marginal product of labor with the wage for each combination of age and skill groups (as in Card and Lemieux 2001).

38. This problem is ignored in standard papers on wage inequality such as Card and Lemieux (2001) and several others in this large literature because of the implicit assumption that technological change is exogenous.

39. We estimate this equation in the generated synthetic control data rather than in the raw data because the synthetic control procedure allows us to construct a good control group for the treatment firms, which we would not be able to replicate by fitting the model directly to the raw data.

40. In the production function, technology is assumed to operate uniformly across age groups: β_j^s and β_j^u do not vary with t . With this assumption, σ_A can be identified in the data using within municipality and skill movements in age-specific supplies as we have shown. Note that our approach can still identify σ_A from the data even if we allow for age-specific technical change, as long as the age-specific technical change is exogenous (i.e. if they evolve over time in the same way between the treatment and control group).

In the second step of the estimation, we use data from both the treated group and the control group to identify the effects of college openings on technology change. It is here that we face the challenge of separately identifying σ_E and endogenous technical change parameters. Based on equation (2), and equipped with the estimated σ_A and aggregate supplies, we can rewrite the model as

$$\log \frac{w_{jt(D)}^s}{w_{jt(D)}^u} + \frac{1}{\sigma_A} \left(\log \frac{S_{jt(D)}}{U_{jt(D)}} - \log \frac{S_{t(D)}}{U_{t(D)}} \right) = \delta_{t(D)} + b_j - \frac{1}{\sigma_E} \log \frac{S_{t(D)}}{U_{t(D)}} + e_{jtD}, \quad (3)$$

where b_j are age-group dummies and $\delta_{t(D)} \equiv \log(\theta_{st(D)}/\theta_{ut(D)})$ represents the relative technology efficiency, which is specific to each year and treatment group. Our parameters of interest are the sequence of $\delta_{t(D)}$ and σ_E . Following much of the literature on this topic, the relative technology efficiency, we model $\delta_{t(D)}$ as a linear trend (although we could relax this assumption), interpreted as skill-biased technical change (Katz and Murphy 1992):⁴¹ $\delta_{t(D)} = \delta_0 t + \delta_1(t \times D) + \delta_2 D$. The trend in technical change is allowed to vary with treatment. δ_0 represents skill-biased technical change in the synthetic control group, whereas δ_1 represents the incremental skill-biased technical change taking place in the treated group. A positive δ_1 implies endogenous technical change.

In this equation, $S_{t(D)}/U_{t(D)}$ could be correlated with e_{jtD} , leading to biased estimates of all the parameters. More importantly, any exogenous variation in skill supplies is no longer a valid exclusion restriction for $S_{t(D)}/U_{t(D)}$ in this equation because it also has a direct effect on technical change, $\delta_{t(D)}$, through the channel we emphasize in this paper.

To address this issue, we proceed with two different identification strategies. One is to use an external estimate of σ_E from the literature to back out the technical change parameters. An advantage of this approach is that we can experiment with a range of plausible values of σ_E to gauge the amount of technical change that is needed to match our data.

A second idea is to make one additional timing assumption in order to identify both σ_E and $\delta_{t(D)}$ from our data. A reasonable possibility is to assume that $\theta_{st(D)}/\theta_{ut(D)}$ does not evolve differentially with D in the years immediately following the reform. Formally, this means that $\theta_{st(0)}/\theta_{ut(0)} = \theta_{st(1)}/\theta_{ut(1)}$ for the first M years after the college's opening (although it may obviously vary with t for reasons unrelated to the reform, such as exogenous skill-biased technical change). Under this assumption, we can use the (first M) years immediately after the reform to identify σ_E in equation (3) for fixed $\theta_{st(D)}/\theta_{ut(D)}$, by relating differences in relative wages to (exogenous)

41. The linear trend specification is used for parsimony. In theory, we would be able to identify a more flexible version of the trend from our previous assumption. In particular, with the assumption that $\theta_{st(0)}/\theta_{ut(0)} = \theta_{st(1)}/\theta_{ut(1)}$ up to the first M years after the college opening, we can identify σ_E and σ_A from exogenous changes in the supply of skills, obtained from contrasts between areas with and without college openings. (Because the trend is assumed to be common across these areas in the first M years after the opening of the college.)

differences in skill shares between reform and non-reform areas. Given σ_E and σ_A , we can use the remaining post-reform years to identify the impact of college openings on $\theta_{st(D)}/\theta_{ut(D)}$, $t > M$.⁴²

Our timing assumption can be motivated by existing models of technology adoption and innovation in the literature. In Online Appendix F, we describe a model of endogenous technology adoption following Acemoglu (2007) and explain how the model can imply that there is a threshold for the supply of skilled labor beyond which technical change takes place, and below which it does not (because when there is a fixed cost of technology adoption, and firms only have an incentive to adopt the new technology when return exceeds cost). This provides a theoretical justification for our empirical identification strategy.⁴³

Empirical evidence on the timing of technology change after the reform provides further guidance on our timing assumption. This is more challenging because our R&D data do not allow us to conduct an event-study analysis to identify the timing of the adoption/innovation (see Section 5 for details). Nevertheless, investment in new equipment is likely to be complementary with high-skilled jobs, which can be regarded as a proxy for technology adoption. Figure 6 shows the effects of the reform on investments in different types of new equipment: (i) machinery and equipment, (ii) machinery and equipment but excluding transport equipment (e.g. cars), and (iii) machinery, equipment, and production facilities, where production facilities include infrastructure for production.⁴⁴ Online Appendix Table A.4 shows the corresponding p -values from permutation tests. For all three measures, we find a positive effect on investments in new equipment in the treated municipalities at the time of the reform, which grows over time, and only becomes statistically different from zero around 10 years after the reform. For the measure including production facilities, there is some weak evidence that the positive effects emerge earlier than for measures involving

42. Another intuitive idea would be to use the older workers in the years immediately after the reform to identify $\log(\theta_{st(1)}/\theta_{ut(1)}) - \log(\theta_{st(0)}/\theta_{ut(0)})$ because they did not experience increases in $S_{jt(D)}/U_{jt(D)}$ until much later. Given $\log(\theta_{st(1)}/\theta_{ut(1)}) - \log(\theta_{st(0)}/\theta_{ut(0)})$, one could potentially use the younger workers to identify σ_E and σ_A . Of course, this intuition is not quite correct because, even if $S_{jt(D)}/U_{jt(D)}$ does not increase for older workers, their wages are still potentially affected by increases in this variable among the young. The case where this works exactly is when $\sigma_E = \sigma_A$. Under this assumption, age-specific relative wages only depend on age-specific relative supplies.

43. Other models, including Acemoglu (1998) and Beaudry and Green (2003) also show that an endogenous change in technology only takes place when the supply of skilled workers increases above a certain threshold. For instance, in Acemoglu (1998), new technologies are invented, and inventors devote more effort in the invention of skill-complementary technologies when there are more skilled workers (because when there are more skilled workers, the market for skill-complementary technologies is larger and the inventor is able to obtain higher profits). Given the fixed cost of technology invention, technology innovation will only take place when there is sufficient number of high-skilled workers in the market. Therefore, there may be periods of increasing supply of skill without a corresponding change in technology.

44. Note that these variables measure the investments in new equipment, but the costs of repairing existing equipment are excluded. Measure (ii) is the narrowest measure for machinery and perhaps the closest to equipment used directly in the production process that we are able to obtain from our plant-level data (we do not have data on ICT equipment in the period under study).

only machinery, which could be due to firms' first upgrading production facilities before investing in new machines. The fact that these impacts are all relatively small (and statistically not different from zero) immediately after the reform provides some justification of our M -year lag identifying assumption, where substantial technology effects may only emerge after M years of the reform.

3.3.3. Can the Model Explain the Movement of Wages and Supply without Endogenous Skill-Biased Technical Change?. Can the model explain the movement of relative wages and supply shifts shown in Figure 1, without any endogenous skill-biased technical change? Suppose that there is no endogenous SBTC, which means that $d(\theta_{st}/\theta_{ut}) = 0$, or equivalently, $\theta_{st(1)}/\theta_{ut(1)} = \theta_{st(0)}/\theta_{ut(0)}$, $\forall t$. Then, take the difference between treated group and the synthetic control within age groups, we get

$$d \left(\log \frac{w_{jt}^s}{w_{jt}^u} \right) = -\frac{1}{\sigma_E} d \left(\log \frac{S_t}{U_t} \right) - \frac{1}{\sigma_A} d \left(\log \frac{S_{jt}}{U_{jt}} - \log \frac{S_t}{U_t} \right), j = \{1, 2\}. \quad (4)$$

For old workers ($j = 2$), let us assume that $d(\log(S_{2t}/U_{2t})) = 0$, which is consistent with our empirical results for the first 10 years following the reform. In these circumstances, in order to get the positive wage effect on older workers as shown in Figure 1, we need that $\sigma_E > \sigma_A$, which is possible but inconsistent with estimates of σ_E and σ_A in the literature (e.g. Card and Lemieux 2001).

Our data also suggest that, for young workers, $d(\log(S_{1t}/U_{1t}) - \log(S_t/U_t))$ and $d(\log(S_t/U_t))$ are both positive.⁴⁵ Therefore, without any endogenous technical change, the model predicts a decrease in the relative wage of young workers because both terms on the RHS of equation (4) are negative for young workers. This may be consistent with the data immediately after the reform, but not with the evidence for the subsequent years.

3.3.4. Estimation Results. We start by presenting results where we fix σ_E and estimate the implied rate of technical change in treatment and control areas. Figure 5 shows the impact of a college opening on the resulting rate of endogenous technical change (the incremental trend in the unexplained relative wages in the treated group), for different values of the elasticity of substitution between skill groups. The less substitutable skilled and unskilled labor is, the higher the rate of endogenous technical change implied by the model. Notice also that even if skilled and unskilled labor are highly substitutable, the implied endogenous technical change is still estimated to be positive and significantly different from zero. In Card and Lemieux (2001), the estimated σ_E is between 2 and 2.5. At these values, the implied impact of endogenous technical change on the relative wage of skilled workers is slightly over 0.005 (0.5%) per year.

45. S_t/U_t is constructed given the estimated σ_A , $\log(\beta_j^s)$, and $\log(\beta_j^u)$, as explained before. Note that if young and old workers are perfect substitutes, then $d(\log(S_{1t}/U_{1t}) - \log(S_t/U_t)) = 0$.




FIGURE 5. Implied endogenous technical change for a range of σ_E . This figure shows the estimated rate of endogenous SBTC (and the associated 95% CI) from Online Appendix equation (E.6) for given values of σ_E .




FIGURE 6. The effects of the reform on investment in new equipment. This figure presents the synthetic control estimates on the investment in new equipment. In each graph, the year of the reform is normalized to period zero. Each data point represents the mean outcome in a given period, relative to the levels in the year of the reform (where the outcome in the reform year is normalized to zero). Each graph reports the weighted average of all treated municipalities and the corresponding synthetic controls, with weight given by the number of plants in the treated municipality.

TABLE 2. Estimates from the relative labor demand regression.

	(1)	(2)	(3)	(4)
Aggregate supply	-0.549 (0.348)	-0.546 (0.341)	-0.541 (0.325)	-0.492 (0.312)
Trend	0.006 (0.010)	0.006 (0.010)	0.006 (0.010)	0.005 (0.009)
Trend \times post	0.008** (0.004)	0.009** (0.004)	0.010** (0.004)	0.011*** (0.004)
Treated	0.015 (0.018)	0.021 (0.016)	0.027* (0.015)	0.029** (0.014)
Post	-0.024 (0.021)	-0.038* (0.020)	-0.057*** (0.021)	-0.075*** (0.022)
Older worker	0.141*** (0.007)	0.141*** (0.007)	0.141*** (0.007)	0.141*** (0.007)
Constant	-0.761 (0.713)	-0.756 (0.698)	-0.745 (0.665)	-0.644 (0.639)
N	72	72	72	72

Note: This table reports the estimates based on equation (3) and Online Appendix equation (E.7). See Section 3.3 and Online Appendix E for details. As independent variables, each regression includes a constant, a time trend (t), a time trend interacted with the post dummy ($t \times P_{t(D)}$), a treatment group indicator (D), a post dummy ($P_{t(D)}$), an age group indicator (b_j), and the relative supply index. In column (1), we set $M = 2$. In columns (2)–(4), we set $M = 3, 4$, and 5 , respectively. * $p < 0.10$, ** $p < 0.05$, and *** $p < 0.01$.

Table 2 presents the estimated parameters resulting from models where we use the second identification strategy, relying on a timing of response assumption. In column (1), we fix $M = 2$ as our baseline specification, which means that we allow for incremental growth in the relative technology efficiency in the treated group beginning two years after the opening of a new college. We estimate the interaction between the linear trend and the treatment dummy to be positive and statistically significant, showing that the reform leads to an incremental increase in the relative demand for high-skilled workers. Holding relative supplies fixed, a college opening in the municipality leads to an increase in the growth rate of the relative skilled wage of 0.8% per year. The implied elasticity of substitution between college and non-college labor is slightly below 2.⁴⁶

In subsequent columns, we show that our results are robust to alternative assumptions about the value of M . In columns (2)–(4), we set $M = 3, 4$, and 5 , respectively. Each of these alternatives gives very similar conclusions to our baseline specification: There is an incremental increase in the relative technology efficiency in treated municipalities following the reform relative to non-treated municipalities. We show in Online Appendix Figure A.10 that the estimates of the baseline model ($M=2$) provide a reasonably good fit to the skill premiums of the treated group and synthetic control group.

46. By comparison, Katz and Murphy (1992) report an estimate of the same parameter equal to 1.4 (using data for both men and women from the US).

In column (1) of Online Appendix Table A.5, we report the estimated parameters from the first step of the estimation, where we estimate σ_A using exogenous variation in relative supply. The year dummies show a pattern of steeply rising relative returns. The estimates imply an elasticity of substitution between young and old workers of about 3, illustrating the importance of considering imperfect substitutability between young and old workers for a given skill group.⁴⁷

Our results are consistent with the model of endogenous technology adoption in Acemoglu (2007).⁴⁸ In Online Appendix F, we describe this model in its simplest case. In that model, firms in each market have access to the same set of factor-augmenting technologies and choose the type of technology they want to adopt together with skilled and unskilled labor inputs. When skill levels are low and the relative cost of adopting the skilled-biased technology is large enough, the initial optimal choice of technology is the least skill-biased one. When the supply of skilled workers in the market increases, skilled wages initially decrease, holding technology choice constant. As the supply of skilled workers keeps increasing, eventually the marginal benefit from technology adoption exceeds the marginal cost, and firms switch the choice of technology to the skill-biased one. It is not difficult to extend this model to a more dynamic framework where increases in the supply of skilled workers lead to endogenous skill-biased technical change, which, in turn, leads to further increases in the supply of skilled workers. As discussed in Acemoglu (2007), this sort of dynamics may lead to a positive relationship between the quantity of skilled input and the wage of skilled workers, or a long-run, upward-sloping demand for skill.⁴⁹

3.3.5. Ruling Out Alternative Explanations. Could there be other reasons, not related to endogenous technical change, that potentially explain the evolution of relative wages and skill supplies as observed in our data? One potential explanation is that, in addition to shifting the local skill compositions, these colleges had a direct impact on innovation. However, in the period that we study, these colleges did not produce

47. By comparison, Card and Lemieux (2001) report a higher estimate of the same parameter in the range of 4–6 (using data for men from the US). Note, however, that the age groups used in their paper are more disaggregated (five-year age bins defined from ages 26–60 years) than ours. It is also plausible to expect low elasticity of substitution between age groups in a labor market where there is little competition between workers of different age groups. As described in Huttunen, Møen, and Salvanes (2011), the seniority rule is an important feature of the Norwegian labor market. This could be one possible explanation for the fact that the elasticity of substitution between age groups (σ_A) that we estimated is at the lower end of the literature (as reported in Online Appendix Table A.8).

48. Beaudry and Green (2003) suggest an alternative model of endogenous technology adoption which we also could have adapted for our setting. However, because their model is not focused explicitly on explaining why technological adoption can respond to changes in skill supplies, we chose to discuss instead the model in Acemoglu (2007) which is focused on understanding that issue.

49. Another implication of our results is that standard estimates of supply effects and skill-biased technical change in the literature may need to be amended. In Online Appendix G, we have estimated a standard Card and Lemieux (2001) type model using national data for Norway. We show that incorporating our estimate of endogenous technical change in the empirical model halves the estimate of σ_E and more than doubles that rate of skill-biased technical change between 1967 and 1990.

patents or engage in R&D activities that may directly affect technological change in the private sector. Online Appendix Figure A.4 shows that the fraction of higher education R&D expenditures taking place in regional colleges in the 1970s was extremely low (1%–3%).⁵⁰ In addition, although there is no reliable patent data dating back to the 1970s, patent data from the 1990s shows that universities are far more important for generating new patents than regional colleges (see Online Appendix Figure A.12), and this is likely to be even more exacerbated in the 1970s. This is in contrast to new universities in other countries, which have been found to create strong innovation spillovers to the private sector.⁵¹ Therefore, in our context, the role of these new colleges was only to provide college education and not research.

Even though these new colleges did not engage in research activities, one might still be concerned that the opening of a college may increase the relative demand for high-skilled workers since many positions at the new colleges might require university-educated workers, especially those with prior experience in the sector. If this is the case, then the widening of the skilled wage gap over time in reform municipalities could be explained by the growth of the new colleges, and their resulting impact on the demand for skilled workers. We investigated this hypothesis by documenting, using decennial census data, the relative employment size of the college sector in the municipalities with a new college. In the 1980 census, only 1.93% of total employment in reform municipalities was in the “Universities and colleges” sector. Among high-skilled workers, only 2.95% worked as college lecturers in these colleges.⁵² Therefore, any direct demand effect from college openings is likely to be very small. It is also unlikely that the construction of the regional colleges caused a significant increase in the demand for other skilled services, given how small they were relative to the overall size of the economy.⁵³

We also examined the inflow of skilled workers into the municipality when a new college is constructed using an event-study. If these new colleges begin to hire many skilled workers, then we would expect to see more skilled workers moving into the reform municipality from elsewhere, especially in the first few years when the college expands. Online Appendix Figure A.13 shows that there is no differential change between the treated and control municipalities in terms of the inflow of skilled workers (as a fraction of the population size). Overall, we believe that the employment size of these colleges is too small to shift the local labor demand.

50. In the 1990s (outside our period of study), some of these colleges were converted into universities, but the expenditure share in R&D remains small.

51. See, among others, Kantor and Whalley (2014), Toivanen and Väänänen (2016), and Andrews (*forthcoming*).

52. Among all low-skilled workers residing in the reform city, 0.03% worked as college lecturers in 1980. College lecturers are defined using the Norwegian Standard Classification of Occupations 1965 (061–University professors and readers and 062–Other university lecturers).

53. Relatively small colleges are able to produce large increases in the stock of skilled workers because they produce a constant flow of graduates over several years, which increasingly contributes to the existing stock of college-educated workers.

An increase in the supply of young skilled workers could generate an increase in the wages of old skilled workers if they were somewhat complementary in production. This was already discussed above, when we say that, taking for example equation (4), for this results, we would need that $\sigma_E > \sigma_A$ (stronger substitutability between workers of different skill groups than between workers of different ages), which is at odds with the estimates in the literature showing precisely the opposite. Online Appendix Table A.8 reports estimates from some of the central papers in the literature with estimates of σ_E and σ_A . If we take the estimates of Card and Lemieux (2001) or several other papers, keeping technology fixed, then an increase in the supply of skilled young workers leads to a fall in the wages of skilled old workers. Furthermore, as explained above, it is impossible to use such a model (and the absence of endogenous technical change) to also justify an eventual increase in the wages of skilled young workers.

The model that we have presented so far has a single sector in a particular local labor market. A rival model is a trade model, where there are multiple sectors in the local market together with workers who are mobile across different sectors. With many sectors, each with a different skill intensity, an increase in the supply of skilled workers can potentially lead to a rise in output in the skilled-intensive sector and, more generally, a change in the local output mix (across sectors with different skill intensities), with no obvious implications for wages.⁵⁴

In Online Appendix Figure A.14, we compare the shares of total employment, output, and plants in skill-intensive industries between the treatment and the control group. Skill-intensive industries consist of industries with above-median shares of skilled workers in 1967. We do not find large or significant differences between the treated and control groups in terms of output/employment/plant compositions, suggesting that the skilled sector did not expand more than the unskilled sector immediately after the reform.⁵⁵ Therefore, it is unlikely that changes in sectoral composition are substantially affecting our results. Furthermore, in our production function analysis below, we always control for industry fixed effects. This implies that we uncover technical change responses in the firm-level analysis that are identified only from within-industry variation (as opposed to between-industry comparisons).

4. Firm-Level Evidence: Structural Estimates of the Production Function

In Section 3, we used labor market data to show that the relative wages of skilled workers increase following the opening of a college. This suggests that the demand

54. In general, the trade model implies that the *economy-wide* relative labor demand curve is horizontal because relative wage is determined only by output prices and total factor productivity (TFP) (the Rybczynski Theorem).

55. The *p*-values of the differences are reported in Online Appendix Table A.9.

for skilled workers increased simultaneously (and, we argue, endogenously) with the increase in the supply of skilled workers.

In this section, we bolster this conclusion by combining quasi-experimental variation from the college reform with the estimation of production functions for each industry–municipality combination (constructed by aggregating plant-level data by municipality and industry). Instead of using wage data, we directly estimate the productivity of skilled and unskilled workers and the impact of the reform on firm output using firm-level data on inputs and outputs.

Therefore, we do not need to rely on the assumption of competitive labor markets to learn about the impacts of the reform on the productivity of skilled and unskilled labor. This is potentially important because centralized wage bargaining was important in Norway during the period we studied, which means that fluctuations in productivity may not fully pass onto wages (at least in the short-run). The effects of the reform on labor productivity could well be underestimated if only wage data are used.

We show that both the supply of skill and the marginal product of skilled workers increased in the medium run after the introduction of a new college. This occurs in spite of an increase in the amount of skilled labor used by firms. Again, as discussed in Section 3, our conjecture is that these results are driven by endogenous technical change responding to an increase in the abundance of skilled labor.⁵⁶

4.1. Empirical Strategy

As in Section 3.3, we assume that the production function depends on skilled and unskilled labor. In this section, we will also allow it to depend on capital. Specifically, the production function for industry j in municipality c and year t takes the following form:⁵⁷

$$Y_{jct} = F(K_{jct}, S_{jct}, U_{jct}, \theta_{ct}) + \mu_{jct}, \quad (5)$$

where Y_{jct} is the total real value-added output at factor prices (before taxes). The three types of inputs used in production are: K_{jct} , the total real value of capital stock, S_{jct} , the total employment of skilled workers, and U_{jct} , the total employment of unskilled

56. In Online Appendix Figures A.15 and A.16, we report the synthetic control estimates of the effects of the reform on output per worker at the municipality level. The output per worker is measured by taking the mean log value-added per worker over all firms located in the municipality. We find that the reform leads to a large and persistent increase in the value-added output per worker, and the increase is exclusively driven by colleges that produce graduates in STEM fields. Note that an increase in output could be driven by changes in technological efficiency or changes in inputs. The estimated production function would allow us to disentangle and quantify the specific channels that drive this increase in output.

57. Note that the 3-digit industry classification we use is quite detailed, so the cells are narrow. Examples of 3-digit industries include: manufacturing of beverages (SIC 313), textiles (SIC 321), electrical apparatus and supplies (SIC 383), and transport equipment (SIC 384). Altogether, there are 22 industries within manufacturing, and in many cases (more than 10% of the cells), there is only one plant per industry–municipality–year cell. The average number of plants within a municipality–industry–year cell is 6.1, and the median is 3. In our production function analysis, we always use weights by the number of plants within an industry–municipality–year cell. Therefore, the variations we exploit are closest to estimating production function from plant-level data.

workers. Inputs, especially S_{jct} and U_{jct} , may depend on \mathbf{d}_{ct} , a vector of current and lagged reform indicators, defined by $\mathbf{d}_{ct} = \{d_{ct}^\tau\}_{\tau=0}^R$, where d_{ct}^τ is an indicator function $1(t - R_c = \tau)$, R_c is the year of college reform, and τ is the number of years since the reform. θ_{ct} is a vector of skill-augmenting technology parameters that, as mentioned above, may endogenously depend on (S_{jct}, U_{jct}) , and therefore on \mathbf{d}_{ct} . μ_{jct} is a productivity shock.

When discussing identification of the parameters of this model, and ignoring the distinction between workers of different ages (we come back to this point below), it is useful to rewrite skill-specific labor inputs as a function of total employment (L_{jct}) and skill shares (π_{jct}), where

$$\begin{aligned} S_{jct} &= L_{jct} \times \pi_{jct} \\ U_{jct} &= L_{jct} \times (1 - \pi_{jct}). \end{aligned}$$

Therefore, the production function can be equivalently written as

$$Y_{jct} = F(K_{jct}, L_{jct}, \pi_{jct}, \theta_{ct}) + \mu_{jct}.$$

The advantage of this specification relatively to the one above is that we summarize the skill content of labor in a single variable, π_{jct} . This way, we can distinguish, in our exposition, endogeneity problems due to L_{jct} (quantity of labor employed) and to π_{jct} (skill composition of employment).

Our main objective is to study how the marginal product of skilled and unskilled labor depends on θ , which, in turn, depends on \mathbf{d}_{ct} . We assume that \mathbf{d}_{ct} is exogenous (conditional on covariates, location, and time fixed effects) and therefore, leads to exogenous variation in our main input of interest, π_{jct} , independent of productivity shocks μ_{jct} (and independent of other unobserved input choices). As before, we need to face the problem that π_{jct} has a direct effect on Y_{jct} , but it also has an indirect effect (which is the focus of this paper) through θ_{ct} . To separate these two channels, we use one of the same assumptions discussed above: that θ_{ct} does not vary with π_{jct} in the first M years immediately following the opening of a college.

We use a control-function approach to account for the endogeneity of skill shares, using the following econometric model:

$$\begin{aligned} Y_{jct} &= F(K_{jct}, L_{jct}, \pi_{jct}, \mathbf{d}_{ct}^M, \boldsymbol{\beta}) + \mu_{jct} \\ \pi_{jct} &= G(K_{jct}, L_{jct}, \mathbf{d}_{ct}, \boldsymbol{\theta}) + v_{jct}, \end{aligned}$$

where $\mathbf{d}_{ct}^M = \{d_{ct}^\tau\}_{\tau=M}^R$. The excluded instruments for π_{jct} are $\mathbf{d}_{ct} \setminus \mathbf{d}_{ct}^M = \{d_{ct}^\tau\}_{\tau=0}^{M-1}$. We approximate the control function with the following form (analogous to a simple series expansion):

$$E(\mu_{jct}|v_{jct}, K_{jct}, L_{jct}, \mathbf{d}_{ct}) = E(\mu_{jct}|v_{jct}) = \rho_1 v_{jct} + \rho_2 v_{jct}^2. \quad (6)$$

If K_{jct} and L_{jct} could be assumed to be exogenous, then we could implement the following estimator:

$$Y_{jct} = F(K_{jct}, L_{jct}, \pi_{jct}, d_{ct}^M, \beta) + \rho_1 \hat{v}_{jct} + \rho_2 \hat{v}_{jct}^2 + \omega_{jct},$$

where \hat{v}_{jct} is the estimated residual from the first-stage equation for π_{jct} , which controls for the endogeneity of skill share in the production function. The residual ω_{jct} has zero conditional mean once \hat{v}_{jct} is controlled for.

K_{jct} and L_{jct} are, however, likely to be endogenous as well. Although our main focus in the paper is on π_{ct} and θ_{ct} , we use an additional control function proposed by Levinsohn and Petrin (2003). Levinsohn and Petrin (2003) use a structural model of an optimizing firm to derive the conditions under which the intermediate input demand function depends on the firm-specific state variables, including productivity shocks and capital. Under the assumption that the intermediate input demand is monotonic in the productivity shock for all capital, the intermediate input demand function can be inverted to yield a second control function for the unobserved productivity shock.⁵⁸ We refer to Online Appendix I for details of this control function approach. Our results are robust to the treatment of endogeneity in these other inputs.⁵⁹

One practical difficulty is that, as discussed in Section 2.2, we have annual measures of skill shares at the level of the municipality, π_{ct} , but not at the level of industry (or firm). We only observe this municipality–industry specific labor inputs from the decennial Census data. Industry-level information on labor inputs is essential because the firm-level data used in this section to estimate production functions is only available for manufacturing (and therefore, we cannot estimate municipality-level production functions). As a result, we will be less ambitious in the specification of equation (5) than we were for equation (1), by ignoring the age distinction in the definition of labor aggregates. This distinction was illuminating when we focused on wages, and we could explore it because we had detailed data on labor inputs and wages of different groups. The estimation of even the simpler specification of the production function is then made possible by our use of an interpolation procedure to construct an annual

58. The control function approach of using intermediate inputs (the Levinsohn and Petrin (LP) approach) to back-out productivity shocks assumes that such productivity shocks are factor neutral (TFP shocks) and invertible. If we have a factor-biased productivity shock (a shock on the coefficients), then using intermediate inputs alone may not be sufficient (Doraszelski and Jaumandreu 2018). This limitation with the LP approach makes the experimental variations from the college opening reform (including the timing assumption) more valuable in terms of addressing the endogeneity of relative labor inputs (which respond to factor-biased productivity shocks).

59. Notice that we estimate production functions in the original dataset, as opposed to data generated by a synthetic control estimator, as in Section 3.3. Although this has the disadvantage of being a less reliable research design for determining the causal impact of the reform than the synthetic control method, we saw above that standard difference-in-difference estimates did not produce substantially different results. Our goal here is not to have a reduced-form estimate of the causal impact of the reform on, say, plant output, but to estimate the trajectories of the marginal products of skilled and unskilled labor. The advantage of this procedure is that it provides a more standard treatment of these objects of interest. We could, however, have used a procedure analogous to that in Section 3.3, but it would have been much more cumbersome, given the additional controls we are using here (we would need to estimate a synthetic control estimator for each covariate) and the use of the correction in Levinsohn and Petrin (2003).

municipality-level series for skilled and unskilled labor inputs in the manufacturing sector, which we describe in Online Appendix H.

In the above discussion on identification, we abstracted from any covariates. However, we need control for a full set of fixed effects for year and municipality because the college expansion reform is plausibly exogenous only conditional on these covariates. Including them in a typical CES specification such as that of equation (1), augmented with capital and other covariates, is not practical. One alternative, following Kmenta (1967), is to linearize the CES aggregate (equation (1)) around $\rho = 0$ using a second-order Taylor expansion. In Section 4.2 and Online Appendix I, we consider other models and find our main results robust to alternative parameterizations, including a more flexible translog production function and a simpler Cobb–Douglas specification.

One potential concern is that there may be some other unobservable (e.g. at the industry level) that affects technology adoption and that is potentially correlated with skill intensity. If that's the case, then the interaction between the reform and labor inputs does not identify the impact of the reform on the marginal product of labor. In Section 4.2, we show that our results are robust to allowing for industry fixed effects in the productivity of skilled/unskilled labor, capturing permanent unobserved heterogeneity in the labor productivity across different industries (but the coefficient measuring the impact of the reform on the productivity of skilled labor does not depend on industry). Our results are robust to additionally allow pre-reform characteristics at the municipality–industry level to affect the labor productivity. Our identifying assumption is that, conditional on the covariates, there is no correlated unobservable in the parameters of the production function.

4.2. Estimation Results

In Table 3, we report the mean predicted differences in output elasticities of skilled and unskilled labor between the treated and control groups. Estimates come from the CES, the translog, and the Cobb–Douglas specifications of the production function, for different years following the reform. There are two columns for each type of production function. In the first one of each pair (columns (1), (3), and (5)), we only account for the endogeneity of skill shares with the control function in equation (6), while in the second (columns (2), (4), and (6)), we also account for the endogeneity of the total employment with the control function in equation (I.7) in the Online Appendix.⁶⁰

The table has two panels. The top panel shows the impact of the reform on the output elasticities of skilled labor 2, 5, 10, and 15 years after the reform, implied by each of the production functions, as well as the impact of the reform on the average annual growth in the output elasticities of skilled labor. The bottom panel provides estimates of the impact of the reform on the levels and growth of the output elasticity of unskilled labor. If we take, for example, the specification in column (1), five and ten years after the opening of the college, then the output elasticity of skilled workers

60. Results from the first-stage control function equation are reported in Online Appendix Table A.11.

TABLE 3. Production function estimates.

Differences between treated and control group	CES		Translog		Cobb–Douglas	
	(1)	(2)	(3)	(4)	(5)	(6)
<i>Output elasticity: skilled labor</i>						
2 years after reform	0.013 (0.035)	−0.001 (0.029)	0.011 (0.038)	−0.002 (0.027)	0.022 (0.030)	−0.002 (0.024)
5 years after reform	0.042 (0.032)	0.019 (0.027)	0.044 (0.041)	0.031 (0.032)	0.044 (0.029)	0.018 (0.023)
10 years after reform	0.090*** (0.031)	0.055* (0.032)	0.100** (0.049)	0.088* (0.052)	0.080*** (0.029)	0.052* (0.030)
15 years after reform	0.138*** (0.037)	0.090** (0.045)	0.156** (0.061)	0.144* (0.077)	0.116*** (0.032)	0.085** (0.042)
Average growth per year	0.010*** (0.003)	0.007** (0.003)	0.011*** (0.003)	0.011** (0.005)	0.007*** (0.002)	0.007** (0.003)
<i>Output elasticity: unskilled labor</i>						
2 years after reform	0.002 (0.034)	0.007 (0.030)	0.036 (0.033)	0.047 (0.033)	−0.007 (0.029)	0.007 (0.025)
5 years after reform	−0.033 (0.030)	−0.017 (0.026)	−0.013 (0.034)	−0.003 (0.024)	−0.034 (0.028)	−0.016 (0.022)
10 years after reform	−0.092*** (0.032)	−0.057* (0.034)	−0.096* (0.056)	−0.087 (0.059)	−0.080*** (0.030)	−0.054* (0.032)
15 years after reform	−0.150*** (0.042)	−0.098* (0.053)	−0.179** (0.085)	−0.172* (0.104)	−0.126*** (0.039)	−0.093* (0.050)
Average growth per year	−0.012*** (0.003)	−0.008* (0.005)	−0.016** (0.007)	−0.016* (0.009)	−0.009*** (0.003)	−0.008* (0.004)
Skill share control function	Yes	Yes	Yes	Yes	Yes	Yes
LP control function	No	Yes	No	Yes	No	Yes

Note: This table reports the predicted differences in output elasticity (by skill group) between the treated group and the control group at years 2, 5, 10, and 15 after the reform (holding labor inputs at the mean). For the estimating production functions, see Online Appendix I. Columns (1), (3), and (5) include the control function for skill compositions, and columns (2), (4), and (6) further control for unobserved productivity shocks via intermediate inputs. Average growth per year refers to the mean annual change in output elasticity starting from 2 years after the reform. The number of observations = 18,441. Standard errors are clustered at the municipality level and given in parentheses. * $p < 0.10$, ** $p < 0.05$, and *** $p < 0.01$.

is 4.2 and 9 percentage points higher than it would have been in the absence of the reform, respectively. The corresponding numbers for unskilled workers are −0.033 and −0.092, five and ten years after the reform, respectively.

We find that the reform leads to an increase in the productivity of skilled labor and a decrease in the productivity of unskilled labor. Moreover, these impacts grow over time, presumably because the stock of skilled workers in the labor market is also growing over time, and firms are adjusting their technology accordingly. As mentioned above, the Cobb–Douglas specification is the one most restricting the substitution patterns across different inputs. However, allowing for more flexible substitution patterns in the remaining columns of the table does not change the magnitude of our estimates.

Overall, our results are consistent with both absolute skill-biased technological change (the productivity of skilled workers increases) and relative skill-biased technological change (the relative productivity between skilled and unskilled workers

increases). We also estimate increases in factor-neutral productivity change due to the reform, but the parameters are imprecisely estimated.

Although we do not directly observe technology adoption from the firm-level data used in this section, Norwegian technology historians indicate that there was increasing use of ICT technology across different industries and improved organizational structure in the period we study, both of which may favor skilled workers at the cost of unskilled workers. (In Section 5, we show impacts of the reform on R&D activities, from a different dataset.) For instance, as in many other countries, automation (using computers) was introduced in the late 1960s, and these technologies started to spread to many industries also in Norway starting in the 1970s and even more so in the 1980s (see Sogner (2002) and Wicken (1994) for overviews of Norway in this period, and for US evidence, see for instance, Jovanovic and Rousseau (2005)). ICT technologies began being developed and spreading to different types of use in the 70s and 80s (for an overview, see Bresnahan (2010)).⁶¹ The early ICT industries became important in the modernization of many sectors such as metal industries, ship-building, fisheries, including land-based processing of fish, telecommunication, and in the late 1970s and especially 1980s, the oil sector.⁶² For many industries the introduction of ICT technology was combined with the use of new form of organization of production and new types of management, leading to new forms of more flexible production.

Specification Checks. To separate supply and technology effects, we assume that there is no endogenous technological response immediately after the reform. To understand how robust our results are to our exclusion restriction, we estimate the Cobb–Douglas production function using different M -years as exclusion restrictions. Online Appendix Table A.12 shows that the production function estimates are not affected by different assumptions of M . If any, then using the first 5 years as exclusion restrictions implies slightly larger skill-biased technical change. Although using a larger value of M leads to a stronger first stage because it leads to larger and more significant changes in supply, we find it reassuring that our main parameter of interest is not sensitive to the specific timing we use as an exclusion restriction.⁶³

61. One of the largest producers of mini computers was located in Norway called Norwegian Data (Sogner 2002). Combined with a group of engineers at the well-established Norwegian Technical University (NTH), they develop computerized automatization systems for many process industries.

62. For instance, one of the industrial processes where computerized automation took place was in the metal producing sector, where the process of producing metals, for instance, aluminum, was controlled by computerized systems (Sogner 2002). Other sectors where these types of systems were introduced were the large hydro power industry producing electricity, the production of weapons, which was also a quite large industry exporting to the international market, the ship-building industries, and shipping equipment industries. An example from the shipping industry is the systems developed to control the engines from the bridge of the ship, as well as radar systems.

63. In order to assess the magnitude and the significance of the effect in the first M years on skill shares (our exclusion restriction), in columns (2)–(5) of Online Appendix Table A.11, we report regressions where we distinguish the effects in the first M years versus years beyond the first M years in the first-stage regression. The coefficients on $D_{ct} \times (1 - P_{t(D)})$ identify the effects of the reform on skill shares in the first M years, where $M = 2, 3, 4, 5$ (corresponding to columns 2–5). As we include additional years as the

TABLE 4. Production function estimates: robustness checks.

Variable	(1)	(2)	(3)	(4)	(5)	(6)
<i>Output elasticity: skilled labor</i>						
Average growth per year	0.005** (0.002)	0.006*** (0.002)	0.007*** (0.002)	0.006*** (0.002)	0.007*** (0.002)	0.006** (0.002)
<i>Output elasticity: unskilled labor</i>						
Average growth per year	-0.007** (0.003)	-0.009*** (0.003)	-0.009*** (0.003)	-0.008*** (0.003)	-0.010** (0.004)	-0.011*** (0.004)
<i>Output elasticity: capital</i>						
Average growth per year					0.001 (0.004)	
<i>Output elasticity: equipment capital</i>						
Average growth per year						0.003 (0.003)

Note: This table reports the predicted differences in the annual growth of output elasticity (starting from 2 years after the reform, by skill group) between the treated group and the control group (holding labor inputs at the mean). The estimating production function takes the form of Cobb–Douglas production function, under different assumptions of $\beta_{1,ct(D)}$ and $\beta_{2,ct(D)}$ (equation (I.2) in Online Appendix I). Relative to the baseline, Column (1) drops fixed year effects in $\beta_{1,ct(D)}$ and $\beta_{2,ct(D)}$. Column (2) includes industry fixed effects in $\beta_{1,ct(D)}$ and $\beta_{2,ct(D)}$. Column (3) adds initial municipality–industry specific characteristics (including the share of skilled workers and employment size relative to total manufacturing employment in the municipality, both measured in 1960). Column (4) includes treatment group indicator variables (including a treatment group indicator and an indicator for implementing the reform early). Column (5) allows for the reform to impact the productivity of capital. Column (6) is the same as column (5), except that we replace total capital with equipment capital in the production function. The number of observations = 18,441. Standard errors are clustered at the municipality level and given in parentheses. * p < 0.10, ** p < 0.05, and *** p < 0.01.

Under an alternative set of assumptions (arguably less credible), we can allow the technological response to occur immediately after the reform. We present estimates of such a model in Online Appendix Table A.13. These alternative assumptions are that: (1) supply shifts are exogenous and (2) there is no endogenous technical change in the control group because the supply shifts are never large enough to justify it. In this case, we can identify the marginal product of skilled and unskilled labor using data from the control group and the impact of the reform on their productivity using the treatment–control comparison.

Using this model, we find that the productivity of both types of labor is unaffected by the reform in the short-run. The estimated coefficients on the short-run reform indicators interacted with skilled and unskilled labor are, both individually and jointly, insignificantly different from zero. Therefore, this evidence provides support for our identifying assumption that there are no endogenous technological responses in the short-run.

In addition, in Table 4, we show estimates of the average annual growth in the output elasticities of skilled and unskilled labor for five other alternative specifications (see the Online Appendix I for details). We focus on the simpler Cobb–Douglas specification, although our results are similar regardless of which of three specifications of the production function we use. Columns (1) through (4) show that our results are

exclusion restriction, the effects on skill share become stronger. For instance, the first 5 years combined lead to a strongly significant effect on skill shares, with a joint F-test of 4.5.

robust to the inclusion of alternative covariates in $\beta_{1,ct(D)}$ and $\beta_{2,ct(D)}$. Column (1) drops year fixed effects in $\beta_{1,ct(D)}$ and $\beta_{2,ct(D)}$. Column (2) includes industry fixed effects, and column (3) includes initial municipality–industry specific characteristics (from the 1960 census) in $\beta_{1,ct(D)}$ and $\beta_{2,ct(D)}$. As discussed previously, results from these two columns are informative as to whether the interaction between the reform and labor inputs captures the impact of the reform on the marginal product of labor or be contaminated by some other unobservable (e.g. at the industry level) that is potentially correlated with skill intensity. Column (4) adds treatment group indicators, including a treatment group indicator and another indicator based on the timing of the treatment. Column (5) allows for the reform to impact the productivity of capital. The coefficients on labor inputs interacted with reform indicators are robust to the above specifications.

4.3. Quantifying the Technology Effects on Wages

To better visualize the quantitative implications of our estimates, in this section, we simulate the predicted marginal products of skilled and unskilled labor in the post-reform years. We compute the ratio of the two (or the difference in logs) and decompose changes in this variable into effects of changes in endogenous technology and effects of changes in supply.

Denote the log of the relative marginal product of labor in treatment group D by $\Delta_{jct(D)} \equiv \log MP_{jct(D)}^s - \log MP_{jct(D)}^u$, where $\log MP_{jct}^s$ and $\log MP_{jct}^u$ are the log of the predicted marginal product of skilled and unskilled labor, respectively (after taking into account the effect of the reform in each year t). In treated municipalities, the predicted change in $\Delta_{jct(D)}$ that is due to the reform in year t (after the year of reform) can be decomposed into technology effects and supply effects. Using the Cobb–Douglas specification in Online Appendix equation (I.2), we develop the following decomposition:⁶⁴

$$\begin{aligned}\Delta_{jct(1)} - \Delta_{jct(0)} &= \underbrace{(\log \beta_{1,ct(1)} - \log \beta_{2,ct(1)}) - (\log \beta_{1,ct(0)} - \log \beta_{2,ct(0)})}_{\text{technology effects}} \\ &\quad - \underbrace{[(\log S_{jct(1)} - \log U_{jct(1)}) - (\log S_{jct(0)} - \log U_{jct(0)})]}_{\text{supply effects}}.\end{aligned}$$

The advantage of using the Cobb–Douglas is that this decomposition can be written as a simple combination of coefficients and inputs.

Figure 7 presents the results of this decomposition. The figure has three lines, representing the technology effect, the supply effect, and the net effect

64. $\log S_{jct(1)} - \log S_{jct(0)}$ and $\log U_{jct(1)} - \log U_{jct(0)}$ are the treatment effects for skilled- and unskilled-labor inputs. We compute these two terms from the estimates of the first-stage regression. Note that, given the construction of the labor inputs discussed previously, the reform affects labor inputs S_{jct} only through local skill share s_{ct} .




FIGURE 7. Predicted relative wages: decomposing the technology and supply effects. This figure reports the predicted relative wages from the production function estimates. See Section 4.3 for details.

($\Delta_{jct(1)} - \Delta_{jct(0)}$ in the expression above). Our estimates suggest that M periods after the reform, the technology effect dominates the supply effect, resulting in an increase in the level and growth of the relative marginal product of skilled labor (while immediately after the reform, before technological upgrading takes place in firms, the marginal product of skilled labor declines).

Notice that the predicted relative productivity increase in $\Delta_{jct(1)} - \Delta_{jct(0)}$ is much larger than the observed increase in the relative wages of skilled workers observed after the opening of a college. For example, 15 years after the reform, our synthetic control estimates suggest that the relative wages of skilled workers are higher in the treatment areas than in the control areas by 10 percentage points, whereas our simulations from the production function estimates suggest that the gap in the marginal product of skilled and unskilled labor grew by 40 percentage points.

As we said above, because centralized wage bargaining is rather strong in Norway, there may be strong deviations from competitive labor market prices, and fluctuations in productivity may not readily translate into fluctuations in wages. Our implied pass-through rate from the marginal product of skilled workers to their wages is about 25%, which is in line with existing estimates from studies done in similar institutional contexts (Margolis and Salvanes 2001; Barth et al. 2012; Akerman, Gaarder, and Mogstad 2015).⁶⁵ We cannot rule out the possibility that the estimation method (synthetic control for wages versus difference-in-difference estimator for productivity) also plays a role in this divergence, but we showed above that our

65. For instance, Akerman, Gaarder, and Mogstad (2015) find that around 20% of the increase in marginal productivity of skilled workers (due to firms upgrading their internet technology) is passed through to skilled wages.

estimates of the impact of the reform on wages were robust to the specific method used.

5. Additional Evidence from R&D Activities

Our results so far show that the opening of a college leads to an increase in the relative demand for skills, and an increase in the productivity of skilled workers. We interpret these findings as evidence of technical change induced by an increase in the supply of skilled workers. In this section, we present direct evidence that college openings induced firms to invest more in R&D activities, which provides further support for the argument put forth in this paper.

Using firm-level R&D data, we investigate whether the college expansion reform led to an increase in R&D activities. The unit of observation in the R&D data is called “bransjeenhets”, which consists of all plants of a firm with their main activity in the same industrial sector. Following Møen (2005), we link the R&D data to our manufacturing plant-level data using a combination of firm and detailed industry identifiers (3-digit). For a firm with a single plant within a single economic activity, we measure R&D activities at the plant level; for a firm with multiple plants within a single economic activity, we calculate average plant-level values by dividing the total R&D activities by the number of plants with the same economic activity code.⁶⁶

As discussed in Section 2.2, there are fewer observations on R&D than either on wages or firm value-added. The first R&D survey was conducted after the first college was established, and subsequent surveys were conducted only every few years (as opposed to annually). Using these data, we adopt the following empirical specification:

$$\log(Y_{ct}) = \theta_c + \gamma_{s(c)t} + a_1 D_{ct} + \varepsilon_{ct}, \quad (7)$$

where Y_{ct} is the sum of all R&D activities among plants located in municipality c and year t ; D_{ct} is an indicator variable taking value one if the municipality c has a college in year t ; θ_c are fixed municipality effects; and $\gamma_{s(c)t}$ are county-year fixed effects. We consider two measures of R&D activities: (i) firm expenditures in R&D activities and (ii) man-years devoted to R&D activities.⁶⁷

Table 5 reports our estimates of the parameters of equation (7). Columns (1)–(3) show the effects of the college expansion reform on the log total costs of performing

66. The R&D data only samples large firms above a certain threshold in selected years after 1970 (Section 2.2). Therefore, a smaller number of observations in our plant-level data (2,579 plant-years) are linked to the R&D data. Of the linked observations, there are 55% of the firm-activity units with one or two plants and 52% of the firm-activity units located in a single municipality.

67. The total costs of R&D include the internal operating costs and the external procurement costs related to R&D activities. This variable is available in all rounds of the R&D survey (1970, 1975, 1977, 1979, 1981, 1983, and 1985). Variables on the man-years in performing R&D activities are only available in 1970, 1975, 1979, and 1985, and, therefore, resulting in a smaller sample.

TABLE 5. College reform and R&D activities.

	Log total costs of R&D			Log R&D man-years		
	(1)	(2)	(3)	(4)	(5)	(6)
D_{ct}	0.767** (0.362)	0.833*** (0.228)	1.299** (0.552)	0.944*** (0.338)	1.246* (0.677)	1.590 (1.115)
N	859	859	859	469	469	469
Covariates	No. S×Y	Baseline	Munic' trend	No. S×Y	Baseline	Munic' trend

Note: Estimates from equation (7) in the text. Coefficient on D_{ct} shows the effect of the reform. Standard errors are clustered at the municipality level. See Section 5 in the text for details. * $p < 0.10$, ** $p < 0.05$, and *** $p < 0.01$.

R&D activities. The baseline model (column (2)) suggests that a college opening increases total R&D expenditures by over 80%. This estimate is robust to the inclusion of a municipality-specific linear trend (column (3)), and replacing county–year fixed effects with year fixed effects (column (1)). Columns (4)–(6) show that the reform also has a large positive impact on the man-years performing R&D activities in firms. Including municipality-specific linear trend makes the estimates less precise (due to the small sample used in this regression), but overall our estimates point to a large positive increase in the man-years performing R&D activities (which more than doubled) in treatment areas (relative to control areas) following the reform.

6. Conclusion

The leading hypothesis explaining the simultaneous increase in the supply of skills and of the skill premium observed in many countries over the last five years is skill-biased technology change. A large literature asks why SBTC was so pronounced during this period, and one possibility is that it could be an endogenous response to an increase in the relative supply of skilled labor. Whereas predictions from the endogenous technical change hypothesis are shown to be consistent with aggregate supply and wage changes over time, there is little evidence that an exogenous increase in the supply of skills can cause additional investments in skill-biased technologies and a resulting increase in the productivity of skilled labor.

In this paper, we examine the consequences of an exogenous increase in the supply of skilled labor in several local labor markets in Norway, resulting from the construction of new colleges in the 1970s. The reform shifted the skill compositions of the affected areas over time: regions with a new college had more rapid growth in the share of skilled workers than a set of comparison areas without a new college. We use spatial and temporal variation in the availability of new colleges across local labor markets as a natural experiment to identify the impact of changes in the local supply of skills on local labor market outcomes. Our empirical analysis draws on several large and

long panel datasets containing rich firm-level information on production structure and individual-level information on demographics, education, employment, and earnings.

We find that local average skilled earnings, both relative to unskilled and in levels, increased as a response to the new college, which is suggestive of a skill-biased demand shift. Results from our relative labor demand regressions also indicate unobserved technology change favoring college workers relative to high-school workers. Drawing from a large panel of manufacturing firms, our production function estimates also suggest that there are endogenous skill-biased technology investments in response to a college opening because the productivity of high-skilled workers increased after the reform (even after accounting for changes in the capital stock). We interpret our findings using existing models of directed technical change, which predict that an abundance of skilled workers may encourage firms to use more skill-complementary technologies. As a result, the demand for skills may increase, leading to an upward-sloping demand curve in the long-run.

References

- Aakvik, A., K. G. Salvanes, and K. Vaage (2010). "Measuring Heterogeneity in the Returns to Education Using an Education Reform." *European Economic Review*, 54, 483–500.
- Aanensen, Thomas (2010). "Lønnsutvikling for ansatte i skoleverket 1959–2008." *Samfunnsøkonomien*, 64, 4–11.
- Abadie, Alberto, Alexis Diamond, and Jens Hainmueller (2010). "Synthetic Control Methods for Comparative Case Studies: Estimating the Effect of California's Tobacco Control Program." *Journal of the American Statistical Association*, 105, 493–505.
- Acemoglu, D. (1998). "Why do New Technologies Complement Skills? Directed Technical Change and Wage Inequality." *The Quarterly Journal of Economics*, 113, 1055.
- Acemoglu, Daron (2007). "Equilibrium Bias of Technology." *Econometrica*, 75, 1371–1409.
- Acemoglu, Daron and Joshua Angrist (2001). "How Large are Human-Capital Externalities? Evidence from Compulsory-Schooling Laws." In *NBER Macroeconomics Annual 2000*, edited by Ben S. Bernanke and Kenneth Rogoff, Vol. 15. MIT Press, pp. 9–74.
- Acemoglu, Daron and Melissa Dell (2010). "Productivity Differences between and within Countries." *American Economic Journal: Macroeconomics*, 2, 169–88.
- Acemoglu, Daron and Amy Finkelstein (2008). "Input and Technology Choices in Regulated Industries: Evidence from the Health Care Sector." *Journal of Political Economy*, 116, 837–880.
- Akerman, Anders, Ingvil Gaarder, and Magne Mogstad (2015). "The Skill Complementarity of Broadband Internet." *The Quarterly Journal of Economics*, 130, 1781–1824.
- Albrecht, James, Gerard J. van den Berg, and Susan Vroman (2009). "The Aggregate Labor Market Effects of the Swedish Knowledge Lift Program." *Review of Economic Dynamics*, 12, 129–146.
- Andrews, Michael (forthcoming). "How Do Institutions of Higher Education Affect Local Invention? Evidence from the Establishment of US Colleges." *American Economic Journal: Economic Policy*.
- Autor, David H., Lawrence F. Katz, and Melissa S. Kearney (2008). "Trends in US Wage Inequality: Revising the Revisionists." *The Review of Economics and Statistics*, 90, 300–323.
- Autor, David H., Lawrence F. Katz, and Alan B. Krueger (1998). "Computing Inequality: Have Computers Changed the Labor Market?" *The Quarterly Journal of Economics*, 113, 1169–1213.
- Barth, Erling, Bernt Bratsberg, Torbjørn Hægeland, and Oddbjørn Raaum (2012). "Performance Pay, Union Bargaining and Within-Firm Wage Inequality." *Oxford Bulletin of Economics and Statistics*, 74, 327–362.

- Beaudry, P., M. Doms, and E. Lewis (2010). "Should the Personal Computer be Considered a Technological Revolution? Evidence from US Metropolitan Areas." *Journal of Political Economy*, 118, 988–1036.
- Beaudry, Paul and David A. Green (2003). "Wages and Employment in the United States and Germany: What Explains the Differences?" *American Economic Review*, 93(3), 573–602.
- Berman, Eli, John Bound, and Stephen Machin (1998). "Implications of Skill-Biased Technological Change: International Evidence." *Quarterly Journal of Economics*, 113, 1245–1279.
- Bhuller, Manudeep, Magne Mogstad, and Kjell G. Salvanes (2017). "Life-Cycle Earnings, Education Premiums, and Internal Rates of Return." *Journal of Labor Economics*, 35, 993–1030.
- Black, Sandra E., Paul J. Devereux, and Kjell G. Salvanes (2005). "Why the Apple Doesn't Fall Far: Understanding Intergenerational Transmission of Human Capital." *American Economic Review*, 95(1), 437–449.
- Bloom, Nicholas, Mirko Draca, and John Van Reenen (2016). "Trade Induced Technical Change? The Impact of Chinese Imports on Innovation, IT and Productivity." *The Review of Economic Studies*, 83, 87–117.
- Blundell, Richard, David A Green, and Wenchao Jin (2022). "The U.K. as a Technological Follower: Higher Education Expansion and the College Wage Premium." *The Review of Economic Studies*, 89, 142–180.
- Bresnahan, Timothy (2010). "General Purpose Technologies." In *Handbook of the Economics of Innovation*, Vol. 2. Elsevier, pp. 761–791.
- Card, David and Thomas Lemieux (2001). "Can Falling Supply Explain the Rising Return to College for Younger Men? A Cohort-Based Analysis." *The Quarterly Journal of Economics*, 116, 705–746.
- Ciccone, Antonio and Giovanni Peri (2006). "Identifying Human-Capital Externalities: Theory with Applications." *The Review of Economic Studies*, 73, 381–412.
- Davis, Donald R. and Jonathan I. Dingel (2019). "A Spatial Knowledge Economy." *American Economic Review*, 109(1), 153–170.
- Dotzschelski, Ulrich and Jordi Jaumandreu (2018). "Measuring the Bias of Technological Change." *Journal of Political Economy*, 126, 1027–1084.
- Dustmann, Christian and Albrecht Glitz (2015). "How do Industries and firms Respond to Changes in Local Labor Supply?" *Journal of Labor Economics*, 33, 711–750.
- Halvorsen, R., R. Jensen, and F. Foy (1991). "Documentation of the Panel Data Base for Manufacturing (in Norwegian)." Working paper, Statistics Norway.
- Hanlon, W. Walker (2015). "Necessity is the Mother of Invention: Input Supplies and Directed Technical Change." *Econometrica*, 83, 67–100.
- Hanson, Gordon H. and Matthew J. Slaughter (2002). "Labor-Market Adjustment in Open Economies: Evidence from US States." *Journal of International Economics*, 57, 3–29.
- Hunt, Jennifer and Marjolaine Gauthier-Loiselle (2010). "How Much Does Immigration Boost Innovation?" *American Economic Journal: Macroeconomics*, 2, 31–56.
- Huttunen, Kristiina, Jarle Møen, and Kjell G. Salvanes (2011). "How Destructive is Creative Destruction? Effects of Job Loss on Job Mobility, Withdrawal and Income." *Journal of the European Economic Association*, 9, 840–870.
- Johnsen, Bodil, Wold (1999). *Fra universitetsvisjon til høyskoleintegrasjon*. Norwegian Academic Press, Kristiansand.
- Johnson, George E. (1997). "Changes in Earnings Inequality: The Role of Demand Shifts." *Journal of Economic Perspectives*, 11(2), 41–54.
- Jovanovic, Boyan and Peter L. Rousseau (2005). "General Purpose Technologies." In *Handbook of Economic Growth*, edited by Steven Durlauf and Philippe Aghion, Vol. 1. Elsevier, pp. 1181–1224.
- Kahn, Lawrence M. (1998). "Collective Bargaining and the Interindustry Wage Structure: International Evidence." *Economica*, 65, 507–534.
- Kantor, Shawn and Alexander Whalley (2014). "Knowledge Spillovers from Research Universities: Evidence from Endowment Value Shocks." *Review of Economics and Statistics*, 96, 171–188.
- Kantor, Shawn and Alexander Whalley (2019). "Research Proximity and Productivity: Long-Term Evidence from Agriculture." *Journal of Political Economy*, 127, 819–854.

- Katz, Lawrence F. and Kevin M. Murphy (1992). "Changes in Relative Wages, 1963–1987: Supply and Demand Factors." *The Quarterly Journal of Economics*, 107, 35–78.
- Kerr, William R. (2013). "US High-Skilled Immigration, Innovation, and Entrepreneurship: Empirical Approaches and Evidence." NBER Working Paper No. 19377.
- Kerr, William R. and William F. Lincoln (2010). "The Supply Side of Innovation: H-1B Visa Reforms and US Ethnic Invention." *Journal of Labor Economics*, 28, 473–508.
- Klette, Tor Jakob and Zvi Griliches (1996). "The Inconsistency of Common Scale Estimators When Output Prices Are Unobserved and Endogenous." *Journal of Applied Econometrics*, 11, 343–361.
- Kmenta, J (1967). "On Estimation of the CES Production Function." *International Economic Review*, 8, 180–189.
- Levinsohn, James and Amil Petrin (2003). "Estimating Production Functions Using Inputs to Control for Unobservables." *The Review of Economic Studies*, 70, 317–341.
- Lewis, Ethan (2011). "Immigration, Skill Mix, and Capital Skill Complementarity." *The Quarterly Journal of Economics*, 126, 1029–1069.
- Machin, Stephen and John Van Reenen (1998). "Technology and Changes in Skill Structure: Evidence from Seven OECD Countries." *The Quarterly Journal of Economics*, 113, 1215–1244.
- Margolis, David N. and Kjell G. Salvanes (2001). "Do Firms Really Share Rents with Their Workers?" IZA Discussion Paper No. 330.
- Møen, Jarle (2005). "Is Mobility of Technical Personnel a Source of R&D Spillovers?" *Journal of Labor Economics*, 23, 81–114.
- Moene, Karl Ove and Michael Wallerstein (1997). "Pay Inequality." *Journal of Labor Economics*, 15, 403–430.
- Moretti, E. (2004a). "Workers' Education, Spillovers, and Productivity: Evidence from Plant-Level Production Functions." *American Economic Review*, 95(3), 656–690.
- Moretti, Enrico (2004b). "Estimating the Social Return to Higher Education: Evidence from Longitudinal and Repeated Cross-Sectional Data." *Journal of Econometrics*, 121, 175–212.
- Moser, Petra, Alessandra Voena, and Fabian Waldinger (2014). "German Jewish Émigrés and US Invention." *American Economic Review*, 104(10), 3222–55.
- Ottosen-committee (1966–1970). Innstilling om videreutdanning for artianere m.v. fra komiteen til å utrede sprøsmål om vidreutdanning for artianere og andre med tilsvarende grunnutdanning.
- Peri, G. (2012). "The Effect Of Immigration On Productivity: Evidence from US States." *The Review of Economics and Statistics*, 94, 348–358.
- Peri, Giovanni, Kevin Shih, and Chad Sparber (2015). "STEM Workers, H-1B Visas, and Productivity in US Cities." *Journal of Labor Economics*, 33, S225–S255.
- Rauch, James E. (1993). "Productivity Gains from Geographic Concentration of Human Capital: Evidence from the Cities." *Journal of Urban Economics*, 34, 380–400.
- Sognen, Knut (2002). *En liten brikke i et stort spill: den norske IT-industrien fra krise til vekst 1975–2000*. Fagbokforl.
- Toivanen, Otto and Lotta Väänänen (2016). "Education and Invention." *Review of Economics and Statistics*, 98, 382–396.
- Wicken, O. (1994). *Elektronikkentrepenørene* (in Norwegian). Ad Noram Forlag, Oslo.

Supplementary Data

Supplementary data are available at [JEEA](#) online.